
Data Programming Course Exercises

May 4, 2016

Andrea Spanò
andrea.spano@quantide.com¹

¹<mailto:andrea.spano@quantide.com>

Contents

1	Introduction	7
2	Data Object	9
2.1	Vectors	9
2.1.1	Exercise 1	9
2.1.2	Exercise 2	10
2.1.3	Exercise 3	10
2.2	Matrices	11
2.2.1	Exercise 1	11
2.2.2	Exercise 2	11
2.2.3	Exercise 3	11
2.2.4	Exercise 4	13
2.3	Lists	14
2.3.1	Exercise 1	14
2.3.2	Exercise 2	15
2.4	Factors	17
2.4.1	Exercise 1	17
2.4.2	Exercise 2	17
2.5	Data Frames	18
2.5.1	Exercise 1	18
3	Data Import	21
3.1	Text Files	21
3.1.1	Exercise 1	21
3.1.2	Exercise 2	22

3.1.3	Exercise 3	23
3.2	Excel Files	23
3.2.1	Exercise 1	23
3.2.2	Exercise 2	23
3.3	Databases	24
3.3.1	Exercise 1	24
3.4	R Data Files	25
3.4.1	Exercise 1	25
3.4.2	Exercise 2	26
4	Data Manipulation with dplyr	27
4.1	Data	27
4.1.1	flights	28
4.2	Select	30
4.2.1	Exercise 1	30
4.2.2	Exercise 2	30
4.2.3	Exercise 3	31
4.3	Filter	31
4.3.1	Exercise 1	31
4.3.2	Exercise 2	32
4.3.3	Exercise 3	32
4.4	Arrange	34
4.4.1	Exercise 1	34
4.4.2	Exercise 2	34
4.4.3	Exercise 3	35
4.5	Mutate	35
4.5.1	Exercise 1	35
4.5.2	Exercise 2	36
4.6	Summarise	37
4.6.1	Exercise 1	37
4.7	Group_by	37
4.7.1	Exercise 1	37
4.7.2	Exercise 2	38
4.8	Chain multiple operations (%>%)	39

<i>CONTENTS</i>	5
4.8.1 Exercise 1	39
4.8.2 Exercise 2	39
4.8.3 Exercise 3	40
5 Data Visualization with ggplot2	41
5.1 Data	41
5.1.1 iris	41
5.1.2 mpg	42
5.2 Scatterplot	44
5.2.1 Exercise 1	44
5.2.2 Exercise 2	45
5.3 Box Plot	46
5.3.1 Exercise 1	46
5.4 Histogram	48
5.4.1 Exercise 1	48
5.5 Lineplot	49
5.5.1 Exercise 1	49
5.5.2 Exercise 2	50
5.6 Bar graph	51
5.6.1 Exercise 1	51
5.6.2 Exercise 2	52
6 Writing R functions	55
6.1 Writing R functions	55
6.1.1 Exercise 1	55
6.1.2 Exercise 2	56

Chapter 1

Introduction

In this document you will find some exercises about these sections:

- *Data Objects*
- *Data Import and Export*
- *Data Manipulation*
- *Data Visualization with ggplot2*
- *Writing R functions*

Chapter 2

Data Object

2.1 Vectors

2.1.1 Exercise 1

- a. Create a vector, named `vec1`, containing the following values:

1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 15, 20, 25, 30, 35, 40, 45, 50, 60, 70, 80, 90

```
vec1 <- c(1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 15, 20, 25, 30, 35, 40, 45, 50, 60, 70, 80, 90)
```

- b. Select the 5-th element of `vec1`.

```
vec1[5]
```

```
## [1] 5
```

- c. Select the first 10 elements of `vec1`.

```
vec1[1:10]
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

- d. Select all the elements of `vec1` apart from the 2nd and the 6th element.

```
vec1[-c(2,6)]
```

```
## [1] 1 3 4 5 7 8 9 10 15 20 25 30 35 40 45 50 60 70 80 90
```

2.1.2 Exercise 2

- a. Generate a vector, named `vec2`, containing the numbers from 1 to 10 and of length 8, using the function `seq()`.

```
vec2 <- seq(from=1, to=10, length.out = 8)
```

- b. Select the values of `vec2` which are greater than 4.

```
vec2[vec2>4] # or y > 4; b[y]
```

```
## [1] 4.857143 6.142857 7.428571 8.714286 10.000000
```

- c. Select the values of `vec2` which are equal or less than 2 or which are equal or greater than 6.

```
vec2[vec2<=2 | vec2>=6]
```

```
## [1] 1.000000 6.142857 7.428571 8.714286 10.000000
```

2.1.3 Exercise 3

- a. Generate the following vector using the function `rep()`:

```
vec3 <- c("one", "two", "one", "two", "one", "two")
```

```
vec3 <- rep(c("one", "two"), times=3)
```

- b. Generate a new vector, named `vec5`, combining the previous vector, `vec3`, with the following one:

```
vec4 <- c("three", "four")
```

```
vec5 <- c(vec3, vec4)
vec5
```

```
## [1] "one" "two" "one" "two" "one" "two" "three" "four"
```

2.2 Matrices

2.2.1 Exercise 1

Generate a matrix, named `mat1`, with 5 rows and 3 columns, using `matrix()` function:

```
##      [,1] [,2] [,3]
## [1,]    1    2    3
## [2,]    4    5    6
## [3,]    7    8    9
## [4,]   10   11   12
## [5,]   13   14   15
```

```
mat1 <- matrix(1:15, nrow = 5, ncol = 3, byrow = TRUE)
```

2.2.2 Exercise 2

Starting from the following vector:

```
mat2 <- 1:8
```

Generate a matrix with 2 rows and 4 columns using `dim()` function.

```
dim(mat2) <- c(2,4)
mat2
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    3    5    7
## [2,]    2    4    6    8
```

2.2.3 Exercise 3

- Generate a matrix, named `mat3`, combining the following columns:

```
a <- 1:3
b <- 7:9
c <- 8:6
```

```
mat3 <- cbind(a,b,c)
mat3
```

```
##      a b c
## [1,] 1 7 8
## [2,] 2 8 7
## [3,] 3 9 6
```

b. Add the following row to `mat3`:

```
d <- 4:6
```

```
mat3 <-rbind(mat3, d)
mat3
```

```
##      a b c
##      1 7 8
##      2 8 7
##      3 9 6
## d 4 5 6
```

2.2.4 Exercise 4

Considering the following matrix, named `mat4`:

```
mat4 <- matrix(1:24, nrow = 6, ncol = 4, byrow = TRUE)
mat4

##      [,1] [,2] [,3] [,4]
## [1,]    1    2    3    4
## [2,]    5    6    7    8
## [3,]    9   10   11   12
## [4,]   13   14   15   16
## [5,]   17   18   19   20
## [6,]   21   22   23   24
```

- a. Select the third and the fifth row of `mat4`.

```
mat4[c(3,5),]

##      [,1] [,2] [,3] [,4]
## [1,]    9   10   11   12
## [2,]   17   18   19   20
```

- b. Select all columns of `mat4` apart from the first.

```
mat4[, -1]

##      [,1] [,2] [,3]
## [1,]    2    3    4
## [2,]    6    7    8
## [3,]   10   11   12
## [4,]   14   15   16
## [5,]   18   19   20
## [6,]   22   23   24
```

- c. Select second and third rows and second and third columns of `mat4`.

```
mat4[2:3, 2:3] # or mat4[c(2,3) , c(2,3)]

##      [,1] [,2]
## [1,]    6    7
## [2,]   10   11
```

2.3 Lists

2.3.1 Exercise 1

- a. Generate a list, named `list1` that contains the following R elements:

```
vec <- 1:10
mat <- matrix(1:9, ncol = 3)
name <- "Oscar"

list1 <- list(vec = 1:10, mat = matrix(1:9, ncol = 3), name = "Oscar")
list1

## $vec
## [1] 1 2 3 4 5 6 7 8 9 10
##
## $mat
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9
##
## $name
## [1] "Oscar"
```

- b. Add to `list1` the following element:

```
letters <- c("a", "b", "c", "d")

list1$letters <- letters
list1

## $vec
## [1] 1 2 3 4 5 6 7 8 9 10
##
## $mat
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9
##
## $name
## [1] "Oscar"
##
## $letters
## [1] "a" "b" "c" "d"
```

2.3.2 Exercise 2

Given the following list, named `list2`:

```
list2 <- list(vec = c(1,3,5,7,8), mat = matrix(1:12, ncol = 4),
             sub_list = list(names = c("Veronica", "Enrico", "Andrea", "Anna"),
                             numbers = 1:4))

list2

## $vec
## [1] 1 3 5 7 8
##
## $mat
##      [,1] [,2] [,3] [,4]
## [1,]    1    4    7   10
## [2,]    2    5    8   11
## [3,]    3    6    9   12
##
## $sub_list
## $sub_list$names
## [1] "Veronica" "Enrico"  "Andrea"  "Anna"
##
## $sub_list$numbers
## [1] 1 2 3 4
```

a. Extract the first element of `list2`.

```
list2[[1]]

## $vec
## [1] 1 3 5 7 8
```

b. Extract the objects contained in the first element of `list2`.

```
list2[[1]]

## [1] 1 3 5 7 8
```

c. Extract the element named `sub_list` of `list2`.

```
list2$sub_list
```

```
## $names
## [1] "Veronica" "Enrico"   "Andrea"   "Anna"
##
## $numbers
## [1] 1 2 3 4
```

d. Extract the second rows of the matrix included in the second element of `list2`.

```
list2[[2]][2,] # or list2$mat[2,]

## [1] 2 5 8 11
```


2.4 Factors

2.4.1 Exercise 1

Starting from the vector:

```
fac1 <- c("F", "F", "M", "M", "F")
```

Generate the corresponding factor with two levels: “F” and “M”

```
fac1 <- factor(fac1, levels = c("F", "M"))
fac1
```

```
## [1] F F M M F
## Levels: F M
```

2.4.2 Exercise 2

Starting from the vector:

```
fac2 <- c(1, 1, 1, 2, 2, 2)
```

- a. Generate the corresponding factor considering that 1 = “Female”, 2 = “Male” e 3 = “Trans”.

```
fac2 <- factor(fac2, levels = c(1,2,3), labels = c("Female", "Male", "Trans"))
fac2
```

```
## [1] Female Female Female Male Male Male
## Levels: Female Male Trans
```

- b. Select the all elements of `fac2` apart from “Male”.

```
fac2[fac2!= "Male"]
```

```
## [1] Female Female Female
## Levels: Female Male Trans
```

2.5 Data Frames

2.5.1 Exercise 1

- a. Generate a data frame, named `df1`, corresponding to:

```
##      id      name class mean
## 1    1      Luca   5A  6.0
## 2    2  Chiara   5A  7.0
## 3    3     Lisa   5A  5.0
## 4    4  Matteo   5A  6.5
## 5    5   Alice   5A  7.5
## 6    6   Marco   5B  4.5
## 7    7 Veronica   5B  9.0
## 8    8   Nicola   5B  8.0
## 9    9    Elena   5B  8.5
## 10  10 Daniele   5B  7.0
```

Remember to maintain character vectors as they are, specifying `stringsAsFactors = FALSE`.

```
df1 <- data.frame(id=1:10,
                  name=c("Luca", "Chiara", "Lisa", "Matteo", "Alice", "Marco",
                        "Veronica", "Nicola", "Elena", "Daniele"),
                  class=c(rep("5A", times=5), rep("5B", times=5)),
                  mean= c(6,7,5,6.5,7.5,4.5, 9, 8, 8.5, 7), stringsAsFactors = FALSE)

df1

##      id      name class mean
## 1    1      Luca   5A  6.0
## 2    2  Chiara   5A  7.0
## 3    3     Lisa   5A  5.0
## 4    4  Matteo   5A  6.5
## 5    5   Alice   5A  7.5
## 6    6   Marco   5B  4.5
## 7    7 Veronica   5B  9.0
## 8    8   Nicola   5B  8.0
## 9    9    Elena   5B  8.5
## 10  10 Daniele   5B  7.0
```

```
# Other solution
id <- 1:10
name <- c("Luca", "Chiara", "Lisa", "Matteo", "Alice", "Marco",
          "Veronica", "Nicola", "Elena", "Daniele")
class <- c(rep("5A", times=5), rep("5B", times=5))
mean <- c(6,7,5,6.5,7.5,4.5, 9, 8, 8.5, 7)
df1 <- data.frame(id, name, class, mean, stringsAsFactors = FALSE)
df1
```

```
##      id      name class mean
## 1    1      Luca   5A  6.0
## 2    2    Chiara   5A  7.0
## 3    3      Lisa   5A  5.0
## 4    4    Matteo   5A  6.5
## 5    5     Alice   5A  7.5
## 6    6     Marco   5B  4.5
## 7    7  Veronica   5B  9.0
## 8    8     Nicola   5B  8.0
## 9    9      Elena   5B  8.5
## 10  10  Daniele   5B  7.0
```

b. Select the first 3 rows of `df1`.

```
df1[1:3,]
```

```
##      id      name class mean
## 1    1      Luca   5A     6
## 2    2    Chiara   5A     7
## 3    3      Lisa   5A     5
```

c. Select the last 6 rows and the first 3 columns of `df1`.

```
df1[5:10, 1:3]
```

```
##      id      name class
## 5     5     Alice   5A
## 6     6     Marco   5B
## 7     7  Veronica   5B
## 8     8     Nicola   5B
## 9     9      Elena   5B
## 10  10  Daniele   5B
```

d. Select the column `class` of `df1`.

```
df1$class
```

```
## [1] "5A" "5A" "5A" "5A" "5A" "5B" "5B" "5B" "5B" "5B"
```

e. Convert the column `class` of `df1` in a factor with levels: “5A” and “5B”

```
df1$class <- factor(df1$class, levels = c("5A", "5B"))
df1$class
```

```
## [1] 5A 5A 5A 5A 5A 5B 5B 5B 5B 5B
## Levels: 5A 5B
```

f. How many columns and rows `df1` has?

```
dim(df1) # or ncol(df1) and nrow(df1)

## [1] 10 4
```

g. Generate another dataframe, named `df2` composed by the columns `name` and `mean` of `df1`, specifying the argument `stringsAsFactors = FALSE`.

```
df2 <- data.frame(name = df1$name, mean=df1$mean, stringsAsFactors = FALSE)
df2

##      name mean
## 1    Luca  6.0
## 2  Chiara  7.0
## 3    Lisa  5.0
## 4   Matteo  6.5
## 5    Alice  7.5
## 6    Marco  4.5
## 7  Veronica  9.0
## 8    Nicola  8.0
## 9     Elena  8.5
## 10 Daniele  7.0
```

h. Show the first rows and the structure of `df2`.

```
head(df2)

##      name mean
## 1    Luca  6.0
## 2  Chiara  7.0
## 3    Lisa  5.0
## 4   Matteo  6.5
## 5    Alice  7.5
## 6    Marco  4.5

str(df2)

## 'data.frame':    10 obs. of  2 variables:
## $ name: chr  "Luca" "Chiara" "Lisa" "Matteo" ...
## $ mean: num  6 7 5 6.5 7.5 4.5 9 8 8.5 7
```

Chapter 3

Data Import

First of all, set your working directory in the *data* folder, using `setwd()` function, like in this example

```
setwd("C:/Users/Veronica/Documents/rbase/data")
```

We will work inside this folder.

3.1 Text Files

3.1.1 Exercise 1

- a. Import text file named *"tuscany.txt"* and save it in an R object named `tuscany_df`.
Open the text file before importing it to control if the first row contains column names and to control the field and the decimal separator characters. Remember to not import the character columns as factors.

```
tuscany_df <- read.table("tuscany.txt", header = TRUE, sep = "|",  
                        dec=".", stringsAsFactors = FALSE)
```

- b. Visualize the first rows of `tuscany_df`

```
head(tuscany_df)
```

##	id	sex	year_of_birth	marital_status	income	house_number
## 1	1	M	1969	married	16101.1	5144.0
## 2	2	M	1962	single	17220.0	6158.0
## 3	3	M	1965	divorcee	28801.9	10078.0
## 4	4	F	1968	single	25964.0	11133.7
## 5	5	M	1975	married	16522.5	5078.0
## 6	6	M	1977	married	18124.0	5115.0

##	city_name	province	provincial_acronym
## 1	Riparbella	Pisa	PI
## 2	Capolona	Arezzo	AR
## 3	Pomarance	Pisa	PI
## 4	Cascina	Pisa	PI
## 5	Quarrata	Pistoia	PT
## 6	Castiglion Fiorentino	Arezzo	AR

3.1.2 Exercise 2

Import 7 rows of the text file named “*solar.txt*” skipping the first two rows. Save it in the object `solar_df`.

Open the text file before importing it to control if the first row contains column names and to control the field and the decimal separator characters. Remember to not import the character columns as factors.

```
solar_df <- read.table("solar.txt", header = FALSE, sep = ",",
                      dec=".", stringsAsFactors = FALSE,
                      nrows = 7, skip = 2)

solar_df
```

##	V1	V2	V3	V4
## 1	mar	23877	24671	22455
## 2	apr	24377	23677	23670
## 3	mag	24581	25476	24999
## 4	giu	22154	21998	22451
## 5	lug	20924	21645	23871
## 6	ago	23183	22576	23556
## 7	set	27446	27695	28664

3.1.3 Exercise 3

Considering the following data frame, named `df`:

```
df <- data.frame(col1=1:4, col2=4:1, col3=c("one", "two", "three", "four"),
                 stringsAsFactors = FALSE)
```

Save it in a .txt file named “*exercise-3.txt*” in *data* folder.

```
write.table(df, file="exercise-3.txt")
```

3.2 Excel Files

3.2.1 Exercise 1

- Import .xlsx file “*flowers.xlsx*” using `XLConnect` function `loadWorkbook()` and save it in a R workbook object named `flowers`.

Remember to load `XLConnect` package, supposing it is already installed.

```
require(XLConnect)
```

```
flowers <- loadWorkbook("flowers.xlsx")
```

- Read *iris* sheet with `readWorksheet()` function and save it in `flower_df` object. Then, visualize its first rows.

```
flowers_df <- readWorksheet(flowers, sheet = 'iris')
head(flowers_df)
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
## 1	5.1	3.5	1.4	0.2	setosa
## 2	4.9	3.0	1.4	0.2	setosa
## 3	4.7	3.2	1.3	0.2	setosa
## 4	4.6	3.1	1.5	0.2	setosa
## 5	5.0	3.6	1.4	0.2	setosa
## 6	5.4	3.9	1.7	0.4	setosa

3.2.2 Exercise 2

- Create a new file xlsx, named “*exercise-2.xlsx*”, and save it in the R worksheet object, named `ex_2`. Use: `loadWorkbook()` and `saveWorkbook()` functions of `XLConnect`.

```
require(XLConnect)
ex_2 <- loadWorkbook(filename = "exercise-2.xlsx", create = TRUE)
saveWorkbook(ex_2)
```

- b. Create a sheet, named `df`, in the R workbook object using `createSheet()` function. Remember to save the changes also in `.xlsx` file (use `saveWorkbook()` function).

```
createSheet(object = ex_2, name = 'df')
saveWorkbook(ex_2)
```

- c. Considering the following data frame, named `numbers_df`:

```
numbers_df <- data.frame(a= 1:4, b=c("one", "two", "three", "four"),
                        stringsAsFactors = FALSE)
numbers_df
```

```
##   a    b
## 1 1  one
## 2 2  two
## 3 3 three
## 4 4  four
```

Add it to `df` sheet of `ex_2` R workbook object, starting from row 3 and from column 2. Use the function `writeWorksheet()`. Remember to save the changes also in `.xlsx` file (use `saveWorkbook()` function).

```
writeWorksheet(object = ex_2, data = numbers_df, sheet = "df", startRow = 3, startCol = 3)
saveWorkbook(ex_2)
```

3.3 Databases

3.3.1 Exercise 1

- a. Connect to “*plant.sqlite*” SQLite database, using `dbConnect()` function of `RSQLite` package. Save the connection in an R object, named `con`. Remember to load `RSQLite` package, supposing it is already installed.

```
require(RSQLite)

con <- dbConnect(RSQLite::SQLite(), "plant.sqlite")
```

- b. See the list of available tables in “*plant.sqlite*” db, using `dbListTables()` function.

```
dbListTables(con)
```



```
## [1] "PlantGrowth"
```

- c. See list of fields in “*PlantGrowth*” table of “*plant.sqlite*” db, using `dbListFields()` function.

```
dbListFields(con, name = "PlantGrowth")
```

```
## [1] "weight" "group"
```

- d. Send query to “*PlantGrowth*” table of “*plant.sqlite*” which select the records with `weight` greater than 5.5.

```
dbGetQuery(con, "SELECT * FROM PlantGrowth WHERE weight >= 5.5")
```

```
##  weight group
## 1   5.58  ctrl
## 2   6.11  ctrl
## 3   5.87  trt1
## 4   6.03  trt1
## 5   6.31  trt2
## 6   5.54  trt2
## 7   5.50  trt2
## 8   6.15  trt2
## 9   5.80  trt2
```

- e. Disconnect from the database, using `dbDisconnect()` function.

```
dbDisconnect(con)
```

```
## [1] TRUE
```

3.4 R Data Files

3.4.1 Exercise 1

Given the following data frame, named `df_rdata`:

```
df_rdata <- data.frame(a=1:20, b=20:1)
```

Save it in *.Rda* format in the file “*df_rdata.Rda*”, using `save()` function.

```
save(df_rdata, file = "df_rdata.Rda")
```

```
## [1] TRUE
```

3.4.2 Exercise 2

Load “*drug.Rda*” file into the environment, using `load()` function.

```
load("drug.Rda")
```

Chapter 4

Data Manipulation with dplyr

Load dplyr package, supposing it is already installed.

```
require(dplyr)
```

4.1 Data

All the following exercises are based on the `nycflights13` data, taken from the `nycflights13` package.

So first of all, install and load this package

```
install.packages("nycflights13")  
require(nycflights13)
```

The `nycflights13` package contains information about all flights that departed from NYC (e.g. EWR, JFK and LGA) in 2013: 336,776 flights in total.

```
ls(pos = "package:nycflights13")  
  
## [1] "airlines" "airports" "flights"  "planes"   "weather"
```

To help understand what causes delays, it includes a number of useful datasets:

- `flights`: information about all flights that departed from NYC
- `weather`: hourly meteorological data for each airport;
- `planes`: construction information about each plane;
- `airports`: airport names and locations;

- **airlines**: translation between two letter carrier codes and names.

Let us explore the features of **flights** datasets, which will be used in the following exercises.

```
data("flights")
```

4.1.1 flights

This dataset contains on-time data for all flights that departed from NYC (i.e. JFK, LGA or EWR) in 2013. The data frame has 16 variables and 336776 observations. The variables are organised as follow:

- Date of departure: **year**, **month**, **day**;
- Departure and arrival times (local tz): **dep_time**, **arr_time**;
- Departure and arrival delays, in minutes: **dep_delay**, **arr_delay** (negative times represent early departures/arrivals);
- Time of departure broken in to hour and minutes: **hour**, **minute**;
- Two letter carrier abbreviation: **carrier**;
- Plane tail number: **tailnum**;
- Flight number: **flight**;
- Origin and destination: **origin**, **dest**;
- Amount of time spent in the air: **air_time**;
- Distance flown: **distance**.

```
dim(flights)
```

```
## [1] 336776      16
```

```
head(flights)
```

```
##   year month day dep_time dep_delay arr_time arr_delay carrier tailnum flight
## 1 2013     1   1     517         2     830         11      UA  N14228  1545
## 2 2013     1   1     533         4     850         20      UA  N24211  1714
## 3 2013     1   1     542         2     923         33      AA  N619AA  1141
## 4 2013     1   1     544        -1    1004        -18      B6  N804JB   725
## 5 2013     1   1     554        -6     812        -25      DL  N668DN   461
## 6 2013     1   1     554        -4     740         12      UA  N39463  1696
##   origin dest air_time distance hour minute
## 1   EWR  IAH      227     1400    5      17
## 2   LGA  IAH      227     1416    5      33
```

```
## 3    JFK  MIA      160      1089      5      42
## 4    JFK  BQN      183      1576      5      44
## 5    LGA  ATL      116       762      5      54
## 6    EWR  ORD      150       719      5      54
```

```
str(flights)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   336776 obs. of  16 variables:
## $ year      : int  2013 2013 2013 2013 2013 2013 2013 2013 2013 2013 ...
## $ month     : int   1  1  1  1  1  1  1  1  1  1 ...
## $ day       : int   1  1  1  1  1  1  1  1  1  1 ...
## $ dep_time  : int  517 533 542 544 554 554 555 557 557 558 ...
## $ dep_delay: num    2  4  2 -1 -6 -4 -5 -3 -3 -2 ...
## $ arr_time  : int  830 850 923 1004 812 740 913 709 838 753 ...
## $ arr_delay: num   11 20 33 -18 -25 12 19 -14 -8 8 ...
## $ carrier   : chr   "UA" "UA" "AA" "B6" ...
## $ tailnum   : chr   "N14228" "N24211" "N619AA" "N804JB" ...
## $ flight    : int  1545 1714 1141 725 461 1696 507 5708 79 301 ...
## $ origin    : chr   "EWR" "LGA" "JFK" "JFK" ...
## $ dest      : chr   "IAH" "IAH" "MIA" "BQN" ...
## $ air_time  : num   227 227 160 183 116 150 158 53 140 138 ...
## $ distance  : num   1400 1416 1089 1576 762 ...
## $ hour      : num    5  5  5  5  5  5  5  5  5 ...
## $ minute    : num   17 33 42 44 54 54 55 57 57 58 ...
```

4.2 Select

4.2.1 Exercise 1

Extract the following information:

- month;
- day;
- air_time;
- distance.

```
select(flights, month, day, air_time, distance)
```

```
## Source: local data frame [336,776 x 4]
##
##   month   day air_time distance
##   (int) (int)   (dbl)   (dbl)
## 1     1     1     227     1400
## 2     1     1     227     1416
## 3     1     1     160     1089
## 4     1     1     183     1576
## 5     1     1     116       762
## 6     1     1     150       719
## 7     1     1     158     1065
## 8     1     1      53       229
## 9     1     1     140       944
## 10    1     1     138       733
## ..    ...    ...     ...     ...
```

```
# flights %>% select(month, day, air_time, distance)
```

4.2.2 Exercise 2

Extract all information about `flights` except hour and minute.

```
select(flights, -c(hour, minute))
```

```
## Source: local data frame [336,776 x 14]
##
##   year month   day dep_time dep_delay arr_time arr_delay carrier tailnum
##   (int) (int) (int)   (int)     (dbl)   (int)     (dbl)   (chr)   (chr)
## 1  2013     1     1     517         2     830         11     UA   N14228
## 2  2013     1     1     533         4     850         20     UA   N24211
```

```
## 3 2013 1 1 542 2 923 33 AA N619AA
## 4 2013 1 1 544 -1 1004 -18 B6 N804JB
## 5 2013 1 1 554 -6 812 -25 DL N668DN
## 6 2013 1 1 554 -4 740 12 UA N39463
## 7 2013 1 1 555 -5 913 19 B6 N516JB
## 8 2013 1 1 557 -3 709 -14 EV N829AS
## 9 2013 1 1 557 -3 838 -8 B6 N593JB
## 10 2013 1 1 558 -2 753 8 AA N3ALAA
## .. ... .. ... .. ... .. ... ..
## Variables not shown: flight (int), origin (chr), dest (chr), air_time (dbl),
## distance (dbl)
```

```
# flights %>% select(-c(hour, minute))
```

4.2.3 Exercise 3

Extract `tailnum` variable and rename it into `tail_num`

```
select(flights, tail_num=tailnum)

## Source: local data frame [336,776 x 1]
##
##   tail_num
##   (chr)
## 1 N14228
## 2 N24211
## 3 N619AA
## 4 N804JB
## 5 N668DN
## 6 N39463
## 7 N516JB
## 8 N829AS
## 9 N593JB
## 10 N3ALAA
## .. ...
```

```
# flights %>% select(tail_num=tailnum)
```

4.3 Filter

4.3.1 Exercise 1

Select all flights which delayed more than 1000 minutes at departure.

```
filter(flights, dep_delay > 1000)

## Source: local data frame [5 x 16]
##
##   year month   day dep_time dep_delay arr_time arr_delay carrier tailnum
##   (int) (int) (int)   (int)    (dbl)   (int)    (dbl)   (chr)   (chr)
## 1  2013     1     9     641     1301   1242     1272     HA   N384HA
## 2  2013     1    10    1121     1126   1239     1109     MQ   N517MQ
## 3  2013     6    15    1432     1137   1607     1127     MQ   N504MQ
## 4  2013     7    22     845     1005   1044      989     MQ   N665MQ
## 5  2013     9    20    1139     1014   1457     1007     AA   N338AA
## Variables not shown: flight (int), origin (chr), dest (chr), air_time (dbl),
##   distance (dbl), hour (dbl), minute (dbl)

# flights %>% filter(dep_delay > 1000)
```

4.3.2 Exercise 2

Select all flights which delayed more than 1000 minutes at departure or at arrival.

```
filter(flights, dep_delay > 1000 | arr_delay >1000)

## Source: local data frame [5 x 16]
##
##   year month   day dep_time dep_delay arr_time arr_delay carrier tailnum
##   (int) (int) (int)   (int)    (dbl)   (int)    (dbl)   (chr)   (chr)
## 1  2013     1     9     641     1301   1242     1272     HA   N384HA
## 2  2013     1    10    1121     1126   1239     1109     MQ   N517MQ
## 3  2013     6    15    1432     1137   1607     1127     MQ   N504MQ
## 4  2013     7    22     845     1005   1044      989     MQ   N665MQ
## 5  2013     9    20    1139     1014   1457     1007     AA   N338AA
## Variables not shown: flight (int), origin (chr), dest (chr), air_time (dbl),
##   distance (dbl), hour (dbl), minute (dbl)

# flights %>% filter(dep_delay > 1000 | arr_delay >1000)
```

4.3.3 Exercise 3

Select all flights which took off from “EWR” and landed in “IAH”.

```
filter(flights, origin == "EWR" & dest == "IAH")
```



```
## Source: local data frame [3,973 x 16]
##
##   year month   day dep_time dep_delay arr_time arr_delay carrier tailnum
##   (int) (int) (int)   (int)    (dbl)   (int)    (dbl)   (chr)   (chr)
## 1  2013     1     1     517         2     830         11     UA   N14228
## 2  2013     1     1     739         0    1104         26     UA   N37408
## 3  2013     1     1     908         0    1228          9     UA   N12216
## 4  2013     1     1    1044        -1    1352          1     UA   N667UA
## 5  2013     1     1    1205         5    1503         -2     UA   N39418
## 6  2013     1     1    1356         6    1659         19     UA   N26906
## 7  2013     1     1    1527        12    1854         44     UA   N69059
## 8  2013     1     1    1620         0    1945         23     UA   N18119
## 9  2013     1     1    1725         5    2045         24     UA   N17122
## 10 2013     1     1    1959        -1    2310          3     UA   N76514
## .. ... ..
## Variables not shown: flight (int), origin (chr), dest (chr), air_time (dbl),
##   distance (dbl), hour (dbl), minute (dbl)

# flights %>% filter(origin == "EWR" & dest == "IAH")
```

4.4 Arrange

4.4.1 Exercise 1

Sort the flights in chronological order.

```
arrange(flights, year, month, day)

## Source: local data frame [336,776 x 16]
##
##   year month   day dep_time dep_delay arr_time arr_delay carrier tailnum
##   (int) (int) (int)   (int)     (dbl)   (int)     (dbl)   (chr)   (chr)
## 1  2013     1     1     517         2     830         11     UA   N14228
## 2  2013     1     1     533         4     850         20     UA   N24211
## 3  2013     1     1     542         2     923         33     AA   N619AA
## 4  2013     1     1     544        -1    1004        -18     B6   N804JB
## 5  2013     1     1     554        -6     812        -25     DL   N668DN
## 6  2013     1     1     554        -4     740         12     UA   N39463
## 7  2013     1     1     555        -5     913         19     B6   N516JB
## 8  2013     1     1     557        -3     709        -14     EV   N829AS
## 9  2013     1     1     557        -3     838         -8     B6   N593JB
## 10 2013     1     1     558        -2     753          8     AA   N3ALAA
## .. ... ..
## Variables not shown: flight (int), origin (chr), dest (chr), air_time (dbl),
##   distance (dbl), hour (dbl), minute (dbl)

# flights %>% arrange(year, month, day)
```

4.4.2 Exercise 2

Sort the flights by decreasing arrival delay.

```
arrange(flights, desc(arr_delay))

## Source: local data frame [336,776 x 16]
##
##   year month   day dep_time dep_delay arr_time arr_delay carrier tailnum
##   (int) (int) (int)   (int)     (dbl)   (int)     (dbl)   (chr)   (chr)
## 1  2013     1     9     641    1301    1242    1272     HA   N384HA
## 2  2013     6    15    1432    1137    1607    1127     MQ   N504MQ
## 3  2013     1    10    1121    1126    1239    1109     MQ   N517MQ
## 4  2013     9    20    1139    1014    1457    1007     AA   N338AA
## 5  2013     7    22     845    1005    1044     989     MQ   N665MQ
## 6  2013     4    10    1100     960    1342     931     DL   N959DL
## 7  2013     3    17    2321     911     135     915     DL   N927DA
```

```
## 8 2013 7 22 2257 898 121 895 DL N6716C
## 9 2013 12 5 756 896 1058 878 AA N5DMAA
## 10 2013 5 3 1133 878 1250 875 MQ N523MQ
## .. ... .. ... .. ... .. ... ..
## Variables not shown: flight (int), origin (chr), dest (chr), air_time (dbl),
## distance (dbl), hour (dbl), minute (dbl)
```

```
# flights %>% arrange(desc(arr_delay))
```

4.4.3 Exercise 3

Sort the flights by origin (in alphabetical order) and decreasing arrival delay.

```
arrange(flights, origin, desc(arr_delay))

## Source: local data frame [336,776 x 16]
##
##   year month   day dep_time dep_delay arr_time arr_delay carrier tailnum
##   (int) (int) (int)   (int)     (dbl)   (int)     (dbl)   (chr)   (chr)
## 1  2013     1    10    1121     1126    1239     1109     MQ   N517MQ
## 2  2013    12     5     756     896    1058     878     AA   N5DMAA
## 3  2013     5     3    1133     878    1250     875     MQ   N523MQ
## 4  2013    12    19     734     849    1046     847     DL   N375NC
## 5  2013    12    17     705     845    1026     846     AA   N5EMAA
## 6  2013    11     3     603     798     829     796     DL   N990AT
## 7  2013     2    24    1921     786    2135     773     DL   N348NW
## 8  2013    10    14    2042     702    2255     688     DL   N943DL
## 9  2013     7    21    1555     580    1955     645     AA   N3EMAA
## 10 2013     7     7    2123     653     17     632     VX   N521VA
## .. ... .. ... .. ... .. ... ..
## Variables not shown: flight (int), origin (chr), dest (chr), air_time (dbl),
## distance (dbl), hour (dbl), minute (dbl)
```

```
# flights %>% arrange(origin, desc(arr_delay))
```

4.5 Mutate

4.5.1 Exercise 1

Add the following new variable to the `flights` dataset:

- the speed in miles per hour, named `speed` (`distance / air_time * 60`).

Consider that times are in minutes and distances are in miles.

```
mutate(flights, speed = distance / air_time * 60)

## Source: local data frame [336,776 x 17]
##
##   year month   day dep_time dep_delay arr_time arr_delay carrier tailnum
##   (int) (int) (int)   (int)     (dbl)   (int)     (dbl)   (chr)   (chr)
## 1  2013     1     1     517         2     830         11     UA   N14228
## 2  2013     1     1     533         4     850         20     UA   N24211
## 3  2013     1     1     542         2     923         33     AA   N619AA
## 4  2013     1     1     544        -1    1004        -18     B6   N804JB
## 5  2013     1     1     554        -6     812        -25     DL   N668DN
## 6  2013     1     1     554        -4     740         12     UA   N39463
## 7  2013     1     1     555        -5     913         19     B6   N516JB
## 8  2013     1     1     557        -3     709        -14     EV   N829AS
## 9  2013     1     1     557        -3     838         -8     B6   N593JB
## 10 2013     1     1     558        -2     753          8     AA   N3ALAA
## .. ... ..
## Variables not shown: flight (int), origin (chr), dest (chr), air_time (dbl),
##   distance (dbl), hour (dbl), minute (dbl), speed (dbl)

# flights %>% mutate(speed = distance / air_time * 60)
```

4.5.2 Exercise 2

Add the following new variables to the `flights` dataset:

- the gained time in minutes (named `gain`), defined as the difference between delay at departure and delay at arrival;
- the gain time per hours, defined as `gain / (air_time / 60)`

```
mutate(flights, gain = arr_delay - dep_delay,
       gain_per_hour = gain / (air_time / 60))

## Source: local data frame [336,776 x 18]
##
##   year month   day dep_time dep_delay arr_time arr_delay carrier tailnum
##   (int) (int) (int)   (int)     (dbl)   (int)     (dbl)   (chr)   (chr)
## 1  2013     1     1     517         2     830         11     UA   N14228
## 2  2013     1     1     533         4     850         20     UA   N24211
## 3  2013     1     1     542         2     923         33     AA   N619AA
## 4  2013     1     1     544        -1    1004        -18     B6   N804JB
## 5  2013     1     1     554        -6     812        -25     DL   N668DN
## 6  2013     1     1     554        -4     740         12     UA   N39463
```

```
## 7 2013 1 1 555 -5 913 19 B6 N516JB
## 8 2013 1 1 557 -3 709 -14 EV N829AS
## 9 2013 1 1 557 -3 838 -8 B6 N593JB
## 10 2013 1 1 558 -2 753 8 AA N3ALAA
## .. ... .. ... .. ... .. ... ..
## Variables not shown: flight (int), origin (chr), dest (chr), air_time (dbl),
## distance (dbl), hour (dbl), minute (dbl), gain (dbl), gain_per_hour (dbl)

# flights %>% mutate(gain = arr_delay - dep_delay,
# gain_per_hour = gain / (air_time / 60))
```

4.6 Summarise

4.6.1 Exercise 1

Calculate minimum, mean and maximum delay at arrival. Remember to add `na.rm=TRUE` option to all calculations.

```
summarise(flights, min_delay = min(arr_delay, na.rm=TRUE),
           mean_delay = mean(arr_delay, na.rm=TRUE),
           max_delay = max(arr_delay, na.rm=TRUE))

## Source: local data frame [1 x 3]
##
##   min_delay mean_delay max_delay
##   (dbl)      (dbl)      (dbl)
## 1     -86    6.895377    1272

# flights %>% summarise(min_delay = min(arr_delay, na.rm=TRUE),
#   mean_delay = mean(arr_delay, na.rm=TRUE),
#   max_delay = max(arr_delay, na.rm=TRUE))
```

4.7 Group_by

4.7.1 Exercise 1

Calculate number of flights, minimum, mean and maximum delay at departure for flights by month.

Remember to add `na.rm=TRUE` option to all calculations.

```
by_month <- group_by(flights, month)
```

```
summarise(by_month, min_delay = min(dep_delay, na.rm=TRUE),
          mean_delay = mean(dep_delay, na.rm=TRUE),
          max_delay = max(dep_delay, na.rm=TRUE))
```

```
## Source: local data frame [12 x 4]
##
##   month min_delay mean_delay max_delay
##   (int)   (dbl)     (dbl)     (dbl)
## 1     1      -30  10.036665     1301
## 2     2      -33  10.816843      853
## 3     3      -25  13.227076      911
## 4     4      -21  13.938038      960
## 5     5      -24  12.986859      878
## 6     6      -21  20.846332     1137
## 7     7      -22  21.727787     1005
## 8     8      -26  12.611040      520
## 9     9      -24   6.722476     1014
## 10    10      -25   6.243988      702
## 11    11      -32   5.435362      798
## 12    12      -43  16.576688      896
```

```
# flights %>% group_by(month) %>%
#   summarise(min_delay = min(dep_delay, na.rm=TRUE),
#   mean_delay = mean(dep_delay, na.rm=TRUE),
#   max_delay = max(dep_delay, na.rm=TRUE))
```

4.7.2 Exercise 2

Calculate number of flights (using `n()` operator), mean delay at departure and arrival for flights by origin.

Remember to add `na.rm=TRUE` option to mean calculations.

```
by_origin <- group_by(flights, origin)

summarise(by_origin, n_flights = n(),
          mean_dep_delay = mean(dep_delay, na.rm=TRUE),
          mean_arr_delay = mean(arr_delay, na.rm=TRUE))
```

```
## Source: local data frame [3 x 4]
##
##   origin n_flights mean_dep_delay mean_arr_delay
##   (chr)   (int)     (dbl)         (dbl)
## 1   EWR    120835   15.10795         1109
## 2   JFK    111279   12.11216         1272
## 3   LGA    104662   10.34688          915
```

```
# flights %>% group_by(origin) %>%
#   summarise(n_flights = n(),
#   mean_dep_delay = mean(dep_delay, na.rm=TRUE),
#   mean_arr_delay = max(arr_delay, na.rm=TRUE))
```

4.8 Chain multiple operations (%>%)

4.8.1 Exercise 1

Calculate number of flights, minimum, mean and maximum delay at departure for flights by month.

Remember to add `na.rm=TRUE` option to all calculations.

```
flights %>% group_by(month) %>%
  summarise(min_delay = min(dep_delay, na.rm=TRUE),
  mean_delay = mean(dep_delay, na.rm=TRUE),
  max_delay = max(dep_delay, na.rm=TRUE))
```

```
## Source: local data frame [12 x 4]
##
##   month min_delay mean_delay max_delay
##   (int)   (dbl)      (dbl)      (dbl)
## 1     1      -30  10.036665    1301
## 2     2      -33  10.816843     853
## 3     3      -25  13.227076     911
## 4     4      -21  13.938038     960
## 5     5      -24  12.986859     878
## 6     6      -21  20.846332    1137
## 7     7      -22  21.727787    1005
## 8     8      -26  12.611040     520
## 9     9      -24   6.722476    1014
## 10    10      -25   6.243988     702
## 11    11      -32   5.435362     798
## 12    12      -43  16.576688     896
```

4.8.2 Exercise 2

Calculate the monthly mean gained time in minutes, where the gained time is defined as the difference between delay at departure and delay at arrival. Remember to add `na.rm=TRUE` option to mean calculations.

```
flights %>% group_by(month) %>%
  mutate(gain = dep_delay - arr_delay) %>%
  summarise(mean_gain = mean(gain, na.rm=TRUE))
```

```
## Source: local data frame [12 x 2]
##
##   month mean_gain
##   (int)      (dbl)
## 1     1  3.855519
## 2     2  5.147220
## 3     3  7.356713
## 4     4  2.673124
## 5     5  9.370201
## 6     6  4.244284
## 7     7  4.810872
## 8     8  6.529872
## 9     9 10.648649
## 10    10  6.400238
## 11    11  4.958993
## 12    12  1.611806
```

4.8.3 Exercise 3

For each destination, select all days where the mean delay at arrival is greater than 30 minutes. Remember to add `na.rm=TRUE` option to mean calculations.

```
flights %>% group_by(dest) %>%
  summarise(mean_arr_delay = mean(arr_delay, na.rm=TRUE)) %>%
  filter(mean_arr_delay > 30)

## Source: local data frame [3 x 2]
##
##   dest mean_arr_delay
##   (chr)      (dbl)
## 1  CAE      41.76415
## 2  OKC      30.61905
## 3  TUL      33.65986
```


Chapter 5

Data Visualization with ggplot2

Load `ggplot2` package, supposing it is already installed.

```
require(ggplot2)
```

5.1 Data

5.1.1 iris

Almost all the following exercises are based on the `iris` data, taken from the `datasets` package. It is a base package so it is already installed and loaded.

```
data("iris")
```

This dataset gives the measurements in centimeters of length and width of sepal and petal, respectively, for 50 flowers from each of 3 species of iris. The species are *Iris setosa*, *versicolor*, and *virginica*.

`iris` dataset contains the following variables:

- `Sepal.Length`: length of iris sepal
- `Sepal.Width`: width of iris sepal
- `Petal.Length`: length of iris petal
- `Petal.Width`: width of iris petal
- `Species`: species of iris

```
dim(iris)
```

```
## [1] 150 5
```

```
head(iris)
```

```
## Sepal.Length Sepal.Width Petal.Length Petal.Width Species
## 1          5.1          3.5          1.4          0.2 setosa
## 2          4.9          3.0          1.4          0.2 setosa
## 3          4.7          3.2          1.3          0.2 setosa
## 4          4.6          3.1          1.5          0.2 setosa
## 5          5.0          3.6          1.4          0.2 setosa
## 6          5.4          3.9          1.7          0.4 setosa
```

```
str(iris)
```

```
## 'data.frame': 150 obs. of 5 variables:
## $ Sepal.Length: num 5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
## $ Sepal.Width : num 3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
## $ Petal.Length: num 1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
## $ Petal.Width : num 0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
## $ Species : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...
```

5.1.2 mpg

Some of the exercises are based on `mpg` dataset, taken from the `datasets` package. It is a base package so it is already installed and loaded.

```
data("mpg")
```

This dataset contains the fuel economy data from 1999 and 2008 for 38 popular models of car.

```
dim(mpg)
```

```
## [1] 234 11
```

```
head(mpg)
```

```
## manufacturer model displ year cyl trans drv cty hwy fl class
## 1 audi a4 1.8 1999 4 auto(l5) f 18 29 p compact
## 2 audi a4 1.8 1999 4 manual(m5) f 21 29 p compact
## 3 audi a4 2.0 2008 4 manual(m6) f 20 31 p compact
## 4 audi a4 2.0 2008 4 auto(av) f 21 30 p compact
## 5 audi a4 2.8 1999 6 auto(l5) f 16 26 p compact
## 6 audi a4 2.8 1999 6 manual(m5) f 18 26 p compact
```

```
str(mpg)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   234 obs. of  11 variables:
## $ manufacturer: chr  "audi" "audi" "audi" "audi" ...
## $ model       : chr  "a4" "a4" "a4" "a4" ...
## $ displ       : num  1.8 1.8 2 2 2.8 2.8 3.1 1.8 1.8 2 ...
## $ year        : int  1999 1999 2008 2008 1999 1999 2008 1999 1999 2008 ...
## $ cyl         : int  4 4 4 4 6 6 6 4 4 4 ...
## $ trans       : chr  "auto(l5)" "manual(m5)" "manual(m6)" "auto(av)" ...
## $ drv         : chr  "f" "f" "f" "f" ...
## $ cty         : int  18 21 20 21 16 18 18 18 16 20 ...
## $ hwy         : int  29 29 31 30 26 26 27 26 25 28 ...
## $ fl          : chr  "p" "p" "p" "p" ...
## $ class       : chr  "compact" "compact" "compact" "compact" ...
```

5.2 Scatterplot

5.2.1 Exercise 1

- Generate a scatterplot to analyze the relationship between `Sepal.Width` and `Sepal.Length` variables.
- Set the size of the point as 3 and their colour (colour and fill arguments as “green”).

```
p1 <- ggplot(data = iris, mapping = aes(x=Sepal.Width, y=Sepal.Length)) +  
  geom_point(size=3, colour="green", fill="green")  
p1
```




Figure 5.1:

5.2.2 Exercise 2

- a. Generate a scatterplot to analyze the relationship between `Petal.Width` and `Petal.Length` variables according to iris species, mapped as `colour` aes.

```
p1 <- ggplot(data = iris, mapping = aes(x=Sepal.Width, y=Sepal.Length, colour=Species)) +  
  geom_point()  
p1
```




Figure 5.2:

5.3 Box Plot

5.3.1 Exercise 1

- Build a box plot to compare the differences of sepal width accordingly to the type of iris species.
- Set the fill of boxes as “#00FFFF”, the colour as “#0000FF” and the outlier colours as “red”.
- Add the plot title: “Boxplot of Sepal.Width vs Species”

```
pl <- ggplot(data=iris, aes(x=Species, y=Sepal.Width)) +  
  geom_boxplot(fill="#00FFFF", colour="#0000FF", outlier.colour = "red") +  
  ggtitle("Boxplot of Sepal.Width vs Species")  
pl
```




Figure 5.3:

5.4 Histogram

5.4.1 Exercise 1

- Represent the distribution of sepal length with an histogram.
- Set bins fill as “hotpink” and colour as “deeppink”.
- Set the number of bins as 15.

```
p1 <- ggplot(data=iris, aes(x=Sepal.Length)) +  
  geom_histogram(fill="hotpink", colour="deeppink", bins=15)  
p1
```




Figure 5.4:

5.5 Lineplot

5.5.1 Exercise 1

Let us suppose that the observations on flowers are taken along time, so let us consider the following dataset:

```
require(dplyr)
iris2 <- iris %>% mutate(time=1:150)
```

- Build a line plot to visualize the Sepal.Length along time.

```
ggplot(data = iris2, mapping = aes(y=Sepal.Width, x= time)) + geom_line()
```




Figure 5.5:

5.5.2 Exercise 2

Let us suppose that the observations on flowers are taken along time, so let us consider the following dataset:

```
iris3 <- iris %>% mutate(time=rep(1:50, times=3))
```

- Build a line plot to visualize the `Sepal.Length` along time, according to the `Species`.
- Set `linetype` as “twodash”.

```
ggplot(data = iris3, mapping = aes(y=Sepal.Length, x= time, colour=Species)) +  
  geom_line(linetype=6)
```




Figure 5.6:

5.6 Bar graph

Let us consider mpg dataset.

5.6.1 Exercise 1

- Represent graphically with a bar graph, how many cars there are for each class.
- Represent horizontal bar and set bar width as 0.6

```
p1 <- ggplot(mpg, aes(class)) +  
  coord_flip() +  
  geom_bar(width=0.6)  
p1
```




Figure 5.7:

5.6.2 Exercise 2

- a. Represent graphically with a bar graph, how many cars there are for each class according to manufacturer.

```
p1 <- ggplot(mpg, aes(class, fill=manufacturer)) +  
  geom_bar()  
p1
```




Figure 5.8:

- b. Represent graphically with a bar graph, the distribution of manufacturer for each class.

```
p1 <- ggplot(mpg, aes(class, fill=manufacturer)) +  
  geom_bar(position = "fill")  
p1
```




Figure 5.9:

Chapter 6

Writing R functions

6.1 Writing R functions

6.1.1 Exercise 1

Write a function, named `compute_summary`, which computes: sum, subtraction, multiplication and division of two numbers. The function arguments should be the two numbers, named as: `x` and `y`. The function should return all amounts computed.

```
compute_summary <- function(x, y){  
  sum_op <- x+y  
  sub_op <- x-y  
  mul_op <- x*y  
  div_op <- x/y  
  return(list(sum_op=sum_op, sub_op=sub_op, mul_op=mul_op, div_op=div_op))  
}
```

```
compute_summary(x=4, y=2)
```

```
## $sum_op  
## [1] 6  
##  
## $sub_op  
## [1] 2  
##  
## $mul_op  
## [1] 8  
##  
## $div_op  
## [1] 2
```

```
compute_summary(x=3, y=7)
```

```
## $sum_op
## [1] 10
##
## $sub_op
## [1] -4
##
## $mul_op
## [1] 21
##
## $div_op
## [1] 0.4285714
```

6.1.2 Exercise 2

Write a function, named `compute_gain`, which computes the income by multiplying the amount produced for sale price and then computes the gain by subtracting the costs to income. The function arguments should be: `amount`, `price`, and `costs`; `price` should have a default value equal to 5. The function should return the gain.

```
compute_gain <- function(amount, costs, price=5){
  income = amount * price
  gain = income - costs
  return(gain)
}

compute_gain(amount = 40, costs = 50)

## [1] 150

compute_gain(amount = 100, costs = 70, price = 1)

## [1] 30
```