
Exercises with dplyr and tidyr

March 13, 2017

Enrico Tonini
enrico.tonini@quantide.com¹

¹<mailto:enrico.tonini@quantide.com>

Contents

1	Introduction	7
1.1	Introduction to <code>nycflights13</code> data	7
1.1.1	<code>flights</code>	7
1.1.2	<code>airlines</code>	9
1.1.3	<code>airports</code>	9
1.1.4	<code>planes</code>	10
1.1.5	<code>weather</code>	12
2	Verb functions	17
2.1	<code>select()</code> and its friends	17
2.1.1	Exercise 1	17
2.1.2	Exercise 2	18
2.1.3	Exercise 3	18
2.1.4	Exercise 4	18
2.1.5	Exercise 5	18
2.1.6	Exercise 6	18
2.2	<code>filter()</code> and its friends	18
2.2.1	Exercise 1	19
2.2.2	Exercise 2	19
2.2.3	Exercise 3	20
2.2.4	Exercise 4	20
2.2.5	Exercise 5	21
2.2.6	Exercise 6	21
2.3	<code>arrange()</code>	22
2.3.1	Exercise 1	22

2.3.2	Exercise 2	22
2.3.3	Exercise 3	22
2.4	<code>mutate()</code> and its friends	22
2.4.1	Exercise 1	23
2.4.2	Exercise 2	23
2.4.3	Exercise 3	23
2.4.4	Exercise 4	23
2.4.5	Exercise 5	23
2.4.6	Exercise 6	23
2.5	<code>summarise()</code>	24
2.5.1	Exercise 1	24
2.5.2	Exercise 2	24
2.5.3	Exercise 3	24
3	Grouping data	25
3.1	<code>group_by()</code>	25
3.1.1	Exercise 1	25
3.1.2	Exercise 2	25
3.1.3	Exercise 3	25
3.1.4	Exercise 4	25
3.1.5	Exercise 5	25
3.1.6	Exercise 6	26
4	Do	27
4.1	<code>do</code>	27
4.1.1	Exercise 1	27
4.1.2	Exercise 2	27
4.1.3	Exercise 3	27
4.1.4	Exercise 4	27
5	Combining data	29
5.1	Joins: <code>inner_join()</code> , <code>left_join()</code> , <code>right_join()</code> , etc.	29
5.1.1	Exercise 1	29
5.1.2	Exercise 2	29

<i>CONTENTS</i>	5
5.1.3 Exercise 3	29
5.1.4 Exercise 4	30
5.1.5 Exercise 5	30
5.1.6 Exercise 6	30
6 Tidy data with tidyr	31
6.1 tidyr	31
6.1.1 Exercise 1	31
6.1.2 Exercise 2	31
6.1.3 Exercise 3	32
6.1.4 Exercise 4	32
7 Handling Missing values	33
7.1	33
7.1.1 Exercise 1	33
7.1.2 Exercise 2	33
7.1.3 Exercise 3	33
8 Dates with lubridate	35
8.1 lubridate	35
8.1.1 Exercise 1	35
8.1.2 Exercise 2	36
8.1.3 Exercise 3	36
9 Manipulating strings with stringr	37
9.1	37
9.1.1 Exercise 1	37
9.1.2 Exercise 2	37
9.1.3 Exercise 3	37
10 Case study	39
10.1 Recap exercise	39
10.1.1 Exercise 1	39

Chapter 1

Introduction

In this document you will find some exercises with the `tidyverse` R packages. They are mainly based on the `nycflights13` data, taken from the `nycflights13` package.

1.1 Introduction to `nycflights13` data

The `nycflights13` package contains information about all flights that departed from NYC (e.g. EWR, JFK and LGA) in 2013: 336,776 flights in total.

```
require(nycflights13)
ls(pos = "package:nycflights13")

## [1] "airlines" "airports" "flights"  "planes"   "weather"
```

To help understand what causes delays, it includes a number of useful datasets:

- `flights`: information about all flights that departed from NYC
- `weather`: hourly meteorological data for each airport;
- `planes`: construction information about each plane;
- `airports`: airport names and locations;
- `airlines`: translation between two letter carrier codes and names.

1.1.1 `flights`

This dataset contains on-time data for all flights that departed from NYC (i.e. JFK, LGA or EWR) in 2013. The data frame has 16 variables and 336776 observations. The variables are organised as follow:

- Date of departure: `year`, `month`, `day`;
- Departure and arrival times (local tz): `dep_time`, `arr_time`;
- Departure and arrival delays, in minutes: `dep_delay`, `arr_delay` (negative times represent early departures/arrivals);
- Time of departure broken in to hour and minutes: `hour`, `minute`;
- Two letter carrier abbreviation: `carrier`;
- Plane tail number: `tailnum`;
- Flight number: `flight`;
- Origin and destination: `origin`, `dest`;
- Amount of time spent in the air: `air_time`;
- Distance flown: `distance`.

```
dim(flights)
```

```
## [1] 336776      16
```

```
head(flights)
```

```
##   year month day dep_time dep_delay arr_time arr_delay carrier tailnum flight
## 1 2013     1   1      517         2      830         11      UA  N14228  1545
## 2 2013     1   1      533         4      850         20      UA  N24211  1714
## 3 2013     1   1      542         2      923         33      AA  N619AA  1141
## 4 2013     1   1      544        -1     1004        -18      B6  N804JB   725
## 5 2013     1   1      554        -6      812        -25      DL  N668DN   461
## 6 2013     1   1      554        -4      740         12      UA  N39463  1696
##   origin dest air_time distance hour minute
## 1   EWR  IAH      227     1400     5      17
## 2   LGA  IAH      227     1416     5      33
## 3   JFK  MIA      160     1089     5      42
## 4   JFK  BQN      183     1576     5      44
## 5   LGA  ATL      116       762     5      54
## 6   EWR  ORD      150       719     5      54
```

```
str(flights)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   336776 obs. of  16 variables:
## $ year      : int  2013 2013 2013 2013 2013 2013 2013 2013 2013 2013 2013 ...
## $ month     : int   1 1 1 1 1 1 1 1 1 1 1 ...
## $ day       : int   1 1 1 1 1 1 1 1 1 1 1 ...
## $ dep_time  : int  517 533 542 544 554 554 555 557 557 558 ...
```



```
## $ dep_delay: num  2 4 2 -1 -6 -4 -5 -3 -3 -2 ...
## $ arr_time : int  830 850 923 1004 812 740 913 709 838 753 ...
## $ arr_delay: num  11 20 33 -18 -25 12 19 -14 -8 8 ...
## $ carrier : chr  "UA" "UA" "AA" "B6" ...
## $ tailnum : chr  "N14228" "N24211" "N619AA" "N804JB" ...
## $ flight : int  1545 1714 1141 725 461 1696 507 5708 79 301 ...
## $ origin : chr  "EWR" "LGA" "JFK" "JFK" ...
## $ dest : chr  "IAH" "IAH" "MIA" "BQN" ...
## $ air_time : num  227 227 160 183 116 150 158 53 140 138 ...
## $ distance : num  1400 1416 1089 1576 762 ...
## $ hour : num  5 5 5 5 5 5 5 5 5 ...
## $ minute : num  17 33 42 44 54 54 55 57 57 58 ...
```

1.1.2 airlines

This dataset contains airlines names and their respective carrier codes, it has 2 variables and 16 observations. Data structure shows that both variables involved are categorical.

```
dim(airlines)
```

```
## [1] 16  2
```

```
head(airlines)
```

```
##   carrier      name
## 1     9E   Endeavor Air Inc.
## 2     AA American Airlines Inc.
## 3     AS   Alaska Airlines Inc.
## 4     B6   JetBlue Airways
## 5     DL   Delta Air Lines Inc.
## 6     EV ExpressJet Airlines Inc.
```

```
str(airlines)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   16 obs. of  2 variables:
## $ carrier: Factor w/ 1570 levels "02Q","04Q","05Q",...: 127 143 265 305 485 551 564 584 668 903 ...
## $ name : Factor w/ 1571 levels "40-Mile Air",...: 604 268 236 837 554 635 678 229 751 606 ...
```

1.1.3 airports

This dataset contains useful metadata about airports, that is:

- FAA airport code: `faa`;
- Usual name of the airport: `name`;
- Location of airport: `lat`, `lon`;
- Altitude (in feet): `alt`;
- Timezone offset from GMT: `tz`;
- Daylight savings time zone: `dst` A = Standard US DST: starts on the second Sunday of March, ends on the first Sunday of November U = unknown N = no dst

The data frame has 7 variables and 1397 observations.

```
dim(airports)
```

```
## [1] 1397    7
```

```
head(airports)
```

```
##   faa                name      lat      lon  alt tz dst
## 1 04G      Lansdowne Airport 41.13047 -80.61958 1044 -5  A
## 2 06A  Moton Field Municipal Airport 32.46057 -85.68003  264 -5  A
## 3 06C      Schaumburg Regional 41.98934 -88.10124  801 -6  A
## 4 06N      Randall Airport  41.43191 -74.39156  523 -5  A
## 5 09J      Jekyll Island Airport 31.07447 -81.42778   11 -4  A
## 6 0A9 Elizabethton Municipal Airport 36.37122 -82.17342 1593 -4  A
```

```
str(airports)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   1397 obs. of  7 variables:
##  $ faa : chr  "04G" "06A" "06C" "06N" ...
##  $ name: chr  "Lansdowne Airport" "Moton Field Municipal Airport" "Schaumburg Regional" "Randall Airport" ...
##  $ lat : num  41.1 32.5 42 41.4 31.1 ...
##  $ lon : num  -80.6 -85.7 -88.1 -74.4 -81.4 ...
##  $ alt : int  1044 264 801 523 11 1593 730 492 1000 108 ...
##  $ tz  : num  -5 -5 -6 -5 -4 -4 -5 -5 -5 -8 ...
##  $ dst : chr  "A" "A" "A" "A" ...
```

1.1.4 planes

This dataset contains plane metadata for all plane tailnumbers found in the FAA aircraft registry (American Airways (AA) and Envoy Air (MQ) report fleet numbers rather than tail numbers). The data frame has 9 variables and 3322 observations. The variables are organised as follow:

- Tail number: `tailnum`;
- Year manufactured: `year`;
- Type of plane: `type`;
- Manufacturer and model: `manufacturer`, `model`;
- Number of engines and seats: `engines`, `seats`;
- Average cruising speed in mph: `speed`;
- Type of engine: `engine`.

```
dim(planes)
```

```
## [1] 3322    9
```

```
head(planes)
```

```
##   tailnum year          type      manufacturer      model engines seats
## 1  N10156 2004 Fixed wing multi engine      EMBRAER EMB-145XR      2    55
## 2  N102UW 1998 Fixed wing multi engine AIRBUS INDUSTRIE A320-214      2   182
## 3  N103US 1999 Fixed wing multi engine AIRBUS INDUSTRIE A320-214      2   182
## 4  N104UW 1999 Fixed wing multi engine AIRBUS INDUSTRIE A320-214      2   182
## 5  N10575 2002 Fixed wing multi engine      EMBRAER EMB-145LR      2    55
## 6  N105UW 1999 Fixed wing multi engine AIRBUS INDUSTRIE A320-214      2   182
##   speed      engine
## 1    NA Turbo-fan
## 2    NA Turbo-fan
## 3    NA Turbo-fan
## 4    NA Turbo-fan
## 5    NA Turbo-fan
## 6    NA Turbo-fan
```

```
str(planes)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   3322 obs. of  9 variables:
## $ tailnum      : chr  "N10156" "N102UW" "N103US" "N104UW" ...
## $ year         : int   2004 1998 1999 1999 2002 1999 1999 1999 1999 1999 ...
## $ type         : chr  "Fixed wing multi engine" "Fixed wing multi engine" "Fixed wing multi engine" "Fixed wing multi engine" ...
## $ manufacturer: chr  "EMBRAER" "AIRBUS INDUSTRIE" "AIRBUS INDUSTRIE" "AIRBUS INDUSTRIE" ...
## $ model        : chr  "EMB-145XR" "A320-214" "A320-214" "A320-214" ...
## $ engines      : int   2 2 2 2 2 2 2 2 2 2 ...
## $ seats        : int   55 182 182 182 55 182 182 182 182 182 ...
## $ speed        : int   NA NA NA NA NA NA NA NA NA NA ...
## $ engine       : chr  "Turbo-fan" "Turbo-fan" "Turbo-fan" "Turbo-fan" ...
```

1.1.5 weather

This dataset is about hourly meteorological data for LGA, JFK and EWR. The data frame has 14 variables and 8719 observations. The variables are organised as follow:

- Weather station: `origin` (named `origin` to facilitate merging with `flights` data);
- Time of recording: `year`, `month`, `day`, `hour`;
- Temperature and dewpoint in F: `temp`, `dewp`;
- Relative humidity: `humid`;
- Wind direction (in degrees), speed and gust speed (in mph): `wind_dir`, `wind_speed`, `wind_gust`;
- Precipitation, in inches: `precip`;
- Sea level pressure in millibars: `pressure`;
- Visibility in miles: `visib`.

```
dim(weather)
```

```
## [1] 8719 14
```

```
head(weather)
```

```
##   origin year month day hour  temp  dewp humid wind_dir wind_speed wind_gust
## 1    EWR 2013     1   1    0 37.04 21.92 53.97     230   10.35702  11.91865
## 2    EWR 2013     1   1    1 37.04 21.92 53.97     230   13.80936  15.89154
## 3    EWR 2013     1   1    2 37.94 21.92 52.09     230   12.65858  14.56724
## 4    EWR 2013     1   1    3 37.94 23.00 54.51     230   13.80936  15.89154
## 5    EWR 2013     1   1    4 37.94 24.08 57.04     240   14.96014  17.21583
## 6    EWR 2013     1   1    6 39.02 26.06 59.37     270   10.35702  11.91865
##   precip pressure visib
## 1      0    1013.9    10
## 2      0    1013.0    10
## 3      0    1012.6    10
## 4      0    1012.7    10
## 5      0    1012.8    10
## 6      0    1012.0    10
```

```
str(weather)
```

```
## Classes 'grouped_df', 'tbl_df', 'tbl' and 'data.frame': 8719 obs. of 14 variables:
## $ origin      : chr  "EWR" "EWR" "EWR" "EWR" ...
## $ year        : num  2013 2013 2013 2013 2013 ...
## $ month       : num  1 1 1 1 1 1 1 1 1 1 ...
## $ day         : int  1 1 1 1 1 1 1 1 1 1 ...
## $ hour        : int  0 1 2 3 4 6 7 8 9 10 ...
## $ temp        : num  37 37 37.9 37.9 37.9 ...
## $ dewp        : num  21.9 21.9 21.9 23 24.1 ...
## $ humid       : num  54 54 52.1 54.5 57 ...
## $ wind_dir    : num  230 230 230 230 240 270 250 240 250 260 ...
## $ wind_speed  : num  10.4 13.8 12.7 13.8 15 ...
## $ wind_gust   : num  11.9 15.9 14.6 15.9 17.2 ...
## $ precip      : num  0 0 0 0 0 0 0 0 0 0 ...
## $ pressure    : num  1014 1013 1013 1013 1013 ...
## $ visib       : num  10 10 10 10 10 10 10 10 10 10 ...
## - attr(*, "vars")=List of 3
## ..$ : symbol month
## ..$ : symbol day
## ..$ : symbol hour
## - attr(*, "indices")=List of 8719
## ..$ : int 0
## ..$ : int 1
## ..$ : int 2
## ..$ : int 3
## ..$ : int 4
## ..$ : int 5
## ..$ : int 6
## ..$ : int 7
## ..$ : int 8
## ..$ : int 9
## ..$ : int 10
## ..$ : int 11
## ..$ : int 12
## ..$ : int 13
## ..$ : int 14
## ..$ : int 15
## ..$ : int 16
## ..$ : int 17
## ..$ : int 18
## ..$ : int 19
## ..$ : int 20
## ..$ : int 21
## ..$ : int 22
## ..$ : int 23
## ..$ : int 24
## ..$ : int 25
## ..$ : int 26
## ..$ : int 27
## ..$ : int 28
```

```
## ..$ : int 29
## ..$ : int 30
## ..$ : int 31
## ..$ : int 32
## ..$ : int 33
## ..$ : int 34
## ..$ : int 35
## ..$ : int 36
## ..$ : int 37
## ..$ : int 38
## ..$ : int 39
## ..$ : int 40
## ..$ : int 41
## ..$ : int 42
## ..$ : int 43
## ..$ : int 44
## ..$ : int 45
## ..$ : int 46
## ..$ : int 47
## ..$ : int 48
## ..$ : int 49
## ..$ : int 50
## ..$ : int 51
## ..$ : int 52
## ..$ : int 53
## ..$ : int 54
## ..$ : int 55
## ..$ : int 56
## ..$ : int 57
## ..$ : int 58
## ..$ : int 59
## ..$ : int 60
## ..$ : int 61
## ..$ : int 62
## ..$ : int 63
## ..$ : int 64
## ..$ : int 65
## ..$ : int 66
## ..$ : int 67
## ..$ : int 68
## ..$ : int 69
## ..$ : int 70
## ..$ : int 71
## ..$ : int 72
## ..$ : int 73
## ..$ : int 74
## ..$ : int 75
## ..$ : int 76
## ..$ : int 77
```

```
## ..$ : int 78
## ..$ : int 79
## ..$ : int 80
## ..$ : int 81
## ..$ : int 82
## ..$ : int 83
## ..$ : int 84
## ..$ : int 85
## ..$ : int 86
## ..$ : int 87
## ..$ : int 88
## ..$ : int 89
## ..$ : int 90
## ..$ : int 91
## ..$ : int 92
## ..$ : int 93
## ..$ : int 94
## ..$ : int 95
## ..$ : int 96
## ..$ : int 97
## ..$ : int 98
## .. [list output truncated]
## - attr(*, "group_sizes")= int 1 1 1 1 1 1 1 1 1 ...
## - attr(*, "biggest_group_size")= int 1
## - attr(*, "labels")='data.frame': 8719 obs. of 3 variables:
## ..$ month: num 1 1 1 1 1 1 1 1 1 ...
## ..$ day : int 1 1 1 1 1 1 1 1 1 ...
## ..$ hour : int 0 1 2 3 4 6 7 8 9 10 ...
## ..- attr(*, "vars")=List of 3
## .. ..$ : symbol month
## .. ..$ : symbol day
## .. ..$ : symbol hour
```


Chapter 2

Verb functions

In this section you will find exercises on the basic verbs of data manipulating provided by `dplyr`:

1. `select()`;
2. `filter()`;
3. `arrange()`;
4. `mutate()`;
5. `summarise()`.

```
## Warning in .doLoadActions(where, attach): trying to execute load actions without  
## 'methods' package
```

2.1 `select()` and its friends

Note: all the exercises of this section are based on the `flights` dataset.

```
require(tidyverse)  
require(nycflights13)
```

```
## Loading required package: nycflights13
```

2.1.1 Exercise 1

Extract the following information about flights:

- month;

- day;
- air_time;
- distance.

2.1.2 Exercise 2

Extract all information about flights except hour and minute.

2.1.3 Exercise 3

Select all variables whose name ends in “time”.

2.1.4 Exercise 4

Select all variables whose name contains the word “delay”.

2.1.5 Exercise 5

Select the `tailnum` variable and rename it into `tail_num`.

2.1.6 Exercise 6

Select all the variables and rename the `tailnum` variable into `tail_num`.

```
## Warning in .doLoadActions(where, attach): trying to execute load actions without
## 'methods' package
```

2.2 `filter()` and its friends

Note: all the exercises of this section are based on the `flights` dataset.

```
require(tidyverse)
```

```
## Loading required package: tidyverse
```

```
## Warning in library(package, lib.loc = lib.loc, character.only = TRUE,
## logical.return = TRUE, : there is no package called 'tidyverse'
```

```
require(nycflights13)

## Loading required package: nycflights13
```

2.2.1 Exercise 1

Select all flights which delayed more than 1000 minutes at departure.

```
flights %>% filter(dep_delay > 1000)

## # A tibble: 5 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     1     9     641             900      1301    1242         1530
## 2  2013     1    10    1121            1635      1126    1239         1810
## 3  2013     6    15    1432            1935      1137    1607         2120
## 4  2013     7    22     845            1600      1005    1044         1815
## 5  2013     9    20    1139            1845      1014    1457         2210
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

2.2.2 Exercise 2

Select all flights which delayed more than 1000 minutes at departure or at arrival.

```
flights %>% filter(dep_delay > 1000 | arr_delay > 1000)

## # A tibble: 5 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     1     9     641             900      1301    1242         1530
## 2  2013     1    10    1121            1635      1126    1239         1810
## 3  2013     6    15    1432            1935      1137    1607         2120
## 4  2013     7    22     845            1600      1005    1044         1815
## 5  2013     9    20    1139            1845      1014    1457         2210
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# alternatively
# flights %>% filter(dep_delay > 1000, arr_delay > 1000)
```

2.2.3 Exercise 3

Select all flights which took off from “EWR” and landed in “IAH” on Christmas Day.

```
flights %>% filter(origin == "EWR" & dest == "IAH" & month == 12 & day == 25)
```

```
## # A tibble: 8 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
## 1	2013	12	25	524	515	9	805	814
## 2	2013	12	25	753	747	6	1038	1048
## 3	2013	12	25	1018	1015	3	1310	1316
## 4	2013	12	25	1442	1345	57	1730	1646
## 5	2013	12	25	1530	1529	1	1836	1826
## 6	2013	12	25	1628	1630	-2	1944	1925
## 7	2013	12	25	1843	1804	39	2141	2113
## 8	2013	12	25	2003	2006	-3	2304	2314

```
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# alternatively
# flights %>% filter(origin == "EWR", dest == "IAH", month == 12, day == 25)
```

2.2.4 Exercise 4

Select the first five flights in this dataset.

```
flights %>% slice(1:5)
```

```
## # A tibble: 5 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
## 1	2013	1	1	517	515	2	830	819
## 2	2013	1	1	533	529	4	850	830
## 3	2013	1	1	542	540	2	923	850
## 4	2013	1	1	544	545	-1	1004	1022
## 5	2013	1	1	554	600	-6	812	837

```
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

2.2.5 Exercise 5

Select the last ten flights in this dataset.

```
flights %>% slice((n()-9):n())
```

```
## # A tibble: 10 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
##	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
## 1	2013	9	30	2240	2250	-10	2347	7
## 2	2013	9	30	2241	2246	-5	2345	1
## 3	2013	9	30	2307	2255	12	2359	2358
## 4	2013	9	30	2349	2359	-10	325	350
## 5	2013	9	30	NA	1842	NA	NA	2019
## 6	2013	9	30	NA	1455	NA	NA	1634
## 7	2013	9	30	NA	2200	NA	NA	2312
## 8	2013	9	30	NA	1210	NA	NA	1330
## 9	2013	9	30	NA	1159	NA	NA	1344
## 10	2013	9	30	NA	840	NA	NA	1020

```
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

2.2.6 Exercise 6

Extract information about distance for all flights which delayed more than 1000 minutes at departure.

```
flights %>%
  filter(dep_delay > 1000) %>%
  select(distance)
```

```
## # A tibble: 5 × 1
```

	distance
##	<dbl>
## 1	4983
## 2	719
## 3	483
## 4	589
## 5	2586

```
## Warning in .doLoadActions(where, attach): trying to execute load actions without
## 'methods' package
```

2.3 arrange()

Note: all the exercises of this section are based on the `flights` dataset.

```
require(tidyverse)
require(nycflights13)

## Loading required package: nycflights13
```

2.3.1 Exercise 1

Sort the flights in chronological order.

2.3.2 Exercise 2

Sort the flights by decreasing arrival delay.

2.3.3 Exercise 3

Sort the flights by origin (in alphabetical order) and decreasing arrival delay.

```
## Warning in .doLoadActions(where, attach): trying to execute load actions without
## 'methods' package
```

2.4 mutate() and its friends

Note: all the exercises of this section are based on the `flights` dataset.
Times are in minutes and distances are in miles.

```
require(tidyverse)
require(nycflights13)

## Loading required package: nycflights13
```

2.4.1 Exercise 1

Add the following new variables to the `flights` dataset:

- the gained time in minutes, defined as the difference between delay at departure and delay at arrival;
- the speed in miles per hour (`distance / air_time * 60`).

Show only the following variables: delay at departure, delay at arrival, distance, air time and the two new variables (gained time and speed).

2.4.2 Exercise 2

Redo the previous calculations keeping only the new variables.

2.4.3 Exercise 3

After sorting flights in chronological order, for each flight calculate the difference between its delay at arrival and the delay at arrival of the immediately previous flight. Have R to show only the delay variables (delay at departure, delay at arrival and the new variable).

2.4.4 Exercise 4

For each flight calculate the ‘min ranking’ in terms of delay at arrival.

2.4.5 Exercise 5

For each flight calculate the ‘first ranking’ in terms of delay at arrival.

2.4.6 Exercise 6

Create a variable which indicates if a flight took off on time, i.e. departure delay is more than -4 and less than 4 minutes late.

```
## Warning in .doLoadActions(where, attach): trying to execute load actions without  
## 'methods' package
```

2.5 summarise()

Note: all the exercises of this section are based on the `flights` dataset.

```
require(tidyverse)
require(nycflights13)

## Loading required package: nycflights13
```

2.5.1 Exercise 1

Calculate minimum, mean and maximum delay at arrival.

2.5.2 Exercise 2

Calculate minimum, mean and maximum delay at arrival for flights in January.

2.5.3 Exercise 3

Calculate the number of flights are contained in the dataset

```
## Warning in .doLoadActions(where, attach): trying to execute load actions without
## 'methods' package
```


Chapter 3

Grouping data

3.1 `group_by()`

Note: all the exercises of this section are based on the `flights` dataset.

```
library(tidyverse)
library(nycflights13)
```

3.1.1 Exercise 1

Calculate number of flights, minimum, mean and maximum delay at arrival for flights by month.

3.1.2 Exercise 2

Calculate number of flights, mean delay at departure and arrival for flights by origin.

3.1.3 Exercise 3

Calculate the number of flights that go to each possible destination.

3.1.4 Exercise 4

Calculate the number of flights for each day. Save the result in a data frame called `per_day`.

3.1.5 Exercise 5

By exploiting `per_day`, calculate the number of flights for each month. Save the result in a data frame called `per_month`.

3.1.6 Exercise 6

Calculate the mean daily number of flights per month.

```
## Warning in .doLoadActions(where, attach): trying to execute load actions without  
## 'methods' package
```

Chapter 4

Do

4.1 do

Note: all the exercises of this section are based on the `flights` dataset.

```
library(dplyr)
library(nycflights13)
```

4.1.1 Exercise 1

Calculate quartiles (25-, 50- and 75-percentiles) of delay at arrival per origin. Put all three quartiles in a unique column.

4.1.2 Exercise 2

Redo the previous exercise putting the three quartiles in three different columns (hint: use `summarise()`).

4.1.3 Exercise 3

Calculate mean and standard deviation of delay at arrival per origin. Put both statistics in a unique column.

4.1.4 Exercise 4

Redo the previous exercise putting mean and standard deviation in two different columns (hint: use `summarise()`).

```
## Warning in .doLoadActions(where, attach): trying to execute load actions without
## 'methods' package
```


Chapter 5

Combining data

5.1 Joins: `inner_join()`, `left_join()`, `right_join()`, etc.

Note: all the exercises of this section are based on `flights`, `airlines`, `airports` or `planes` datasets.

```
library(dplyr)
library(nycflights13)
```

5.1.1 Exercise 1

Keep only the following variables of the `flights` dataset: `month`, `day`, `hour`, `origin`, `destination` and `carrier`. Save this dataset in a data frame and call it `flights_red`. Through a proper join command, add the carrier name to `flights_red` (this piece of information is available in `airlines`).

5.1.2 Exercise 2

Through a proper join command, add `name`, `latitude`, `longitude` and `altitude` of the origin airport to `flights_red` (these pieces of information are available in `airports`). Do the same also for the destination airport. (If you are able to, try to keep variables about both origin and destination airports in the same final dataset).

5.1.3 Exercise 3

Through the `inner_join()` function, redo the same for the destination airport but keep only the flights whose information is available in both datasets (`flights` and `airports`).

5.1.4 Exercise 4

Redo the exercise 3 by using `full_join()` instead of `inner_join()`. What is the difference in the result?

5.1.5 Exercise 5

Through the `anti_join()` function, extract all the flights from `flights` whose information about destination airport is not available in `airports`.

5.1.6 Exercise 6

Sort the `planes` dataset by increasing year. Then create two datasets: the first will deal with planes older than 2000; the second will deal with planes of 2000 or newer. Finally create a unique dataset where the first rows will deal with the newest planes, whereas the last rows will deal with the oldest planes.

```
## Warning in .doLoadActions(where, attach): trying to execute load actions without
## 'methods' package
```

Chapter 6

Tidy data with tidyr

6.1 tidyr

```
library(tidyverse)
```

6.1.1 Exercise 1

Consider the following dataset:

```
heartrate_wide <- data.frame(
  name = c("Aldo", "Giovanni", "Giacomo"),
  surname = c("Baglio", "Storti", "Poretti"),
  morning = c(67, 80, 64),
  afternoon = c(56, 90, 50)
)
heartrate_wide
```

	name	surname	morning	afternoon
## 1	Aldo	Baglio	67	56
## 2	Giovanni	Storti	80	90
## 3	Giacomo	Poretti	64	50

It represents the heart rate measured on three patients in the morning and in the afternoon. The dataset is in the wide format: change it to the long format through a proper `tidyr` function. Save the result in a data frame and call it `heartrate_long`.

6.1.2 Exercise 2

Starting from `heartrate_long`, come back to a dataset in a wide format through a proper `tidyr` function. The result should be obviously equal to `heartrate_wide`.

6.1.3 Exercise 3

Consider the dataset `heartrate_wide` and unite name and surname of the patients in a unique column through a proper `tidyr` function. Save the result in a new data frame called `heartrate_united`.

6.1.4 Exercise 4

Starting from `heartrate_united`, come back to a dataset where name and surname are in two different columns through a proper `tidyr` function. The result should be obviously equal to `heartrate_wide`.

```
## Warning in .doLoadActions(where, attach): trying to execute load actions without  
## 'methods' package
```


Chapter 7

Handling Missing values

7.1

```
library(tidyverse)
```

7.1.1 Exercise 1

Consider the following dataset:

```
heartrate <- data.frame(  
  name = c("Aldo", "Giovanni", "Giacomo", "Aldo", "Giovanni", "Giacomo", "Giovanni", "Giacomo"),  
  surname = c("Baglio", "Storti", "Poretti", "Baglio", "Storti", "Poretti", "Storti", "Poretti"),  
  when = c("morning", "morning", "morning", "afternoon", "afternoon", "afternoon", "evening", "evening"),  
  heartrate = c(67, 80, 64, 56, 90, 50, 60, 85)  
)
```

It represents the heart rate measured on three patients in the morning, in the afternoon and in the evening. Make explicit any implicit missing value. How many missing values do you see?

7.1.2 Exercise 2

Import data in the file `marks.Rdta`. Missing values have been recorded as “?”. What’s the percentage of missing values in the data? Replace them with `NA` and drop them.

7.1.3 Exercise 3

Import the data `heartrate_NA.Rdta`. Consider all the missing values you find and replace them using the function `fill()` when possible.

```
## Warning in .doLoadActions(where, attach): trying to execute load actions without  
## 'methods' package
```

Chapter 8

Dates with lubridate

8.1 lubridate

Note: all the exercises of this section are based on the `flights` dataset.

```
require(tidyverse)
require(lubridate)

## Loading required package: lubridate

##
## Attaching package: 'lubridate'

## The following object is masked from 'package:base':
##
##      date

require(nycflights13)

## Loading required package: nycflights13
```

8.1.1 Exercise 1

Using the `flights` data, build the variable `dep_date` based on the variables `year`, `month` and `day`. First use the function `unite()` and then the parsing function `ydm()`. Select only the new variable and save the new data frame called `flights_date`.

8.1.2 Exercise 2

Using the dataset, shift all flights by two months. Save it in a separate data frame called `flights_date_2`.

8.1.3 Exercise 3

Take the new date (2 months ahead) and subtract the original variable `date`. `flights_date_2`.

```
## Warning in .doLoadActions(where, attach): trying to execute load actions without  
## 'methods' package
```

Chapter 9

Manipulating strings with stringr

9.1

```
library(tidyverse)
library(stringr)
```

9.1.1 Exercise 1

Import the data `aire_milano_strings.txt` which is a tab delimited file. Find how China has been codified (notice that the file is in Italian) and manipulate that string as you find more comfortable for you. Save the results in a new tibble.

```
## Parsed with column specification:
## cols(
##   Residenza = col_character(),
##   MotivoIscrizioneEstero = col_character(),
##   Num = col_integer()
## )
```

9.1.2 Exercise 2

Using the data modified in exercise 1, find all the countries whose names contain non-alphanumeric characters. Identify what kind of characters they contain.

9.1.3 Exercise 3

Consider now the column with information on the reason for migrating. Count how many different reasons there are and notice that citizenship was recorded in two slightly different ways: “acquisto cittadinanza” and “per acquisto cittadinanza”. Replace one of them so that they are the same.

```
## Warning in .doLoadActions(where, attach): trying to execute load actions without  
## 'methods' package
```

Chapter 10

Case study

10.1 Recap exercise

In this section you will work on a real data set. Using all the tools provided throughout the course, you will manipulate data for better analysing it. In the data folder you find the following three files:

1. rating_final.csv
2. chefmezcuisine.csv
3. userprofile.csv

These are the files you will work on.

10.1.1 Exercise 1

1. First of all you need to import the three files into **R** using the correct **readr** function. In order to find the correct function and to set the right options, you'd better explore the files by opening them in csv (check which is the separator, if there are column names, etc).
2. In order to understand what you are working on, check how many columns and rows each dataset is composed of, and check what type of variables you are working with. If the variables type has not correctly been parsed, parse it manually. If you find many variables, focus on those that you think may be interesting for understanding different people's tastes (for example age of the users, job, etc).
3. Based on the dataset userprofile, create a new dataset with only relevant variables. Among these, keep the variables: userID, birth_year, budget, marital_status, personality, smoker and activity. If you think there are other relevant variables, you may include them in your dataset as well. Call the new dataset "userprofile_reduced".

4. Focus on the data frame `userprofile_reduced`. By exploring the different values recorded for `budget`, you may notice there are missing values. What are they recorded by? Replace all missing values with `NA`. Do the same for all the variables in `userprofile_reduced`.
5. Note that for all users we have the year of birth but we do not have the age. Replace the year of birth with a variable called `age`.
6. All three data frames are now ready to use. Merge the three datasets so that you keep all rows and columns of `rating` and you add all the variables of `chefmezcuisine.csv` and `userprofile.csv`. Call the new data frame `rating_all`.
7. Find the mean of all rating variables. Group data by `placeID` and then sort the tibble so that places with the highest average rating are at the top. Show id of such places and type of cuisine.
8. Find mean and standard deviation of all rating variables. Do you notice differences with regards to ratings of students as compared to people that are employed? Do you find differences in smokers and non smokers? Do you notice large differences in any other group of users?