
Exercises with dplyr and tidyr

September 5, 2017

Quantide
¹

¹[mailto:](#)

Contents

1	Introduction	7
1.1	Introduction to <code>nycflights13</code> data	7
1.1.1	<code>flights</code>	7
1.1.2	<code>airlines</code>	9
1.1.3	<code>airports</code>	10
1.1.4	<code>planes</code>	11
1.1.5	<code>weather</code>	12
2	Verb functions	15
2.1	<code>select()</code> and its friends	15
2.1.1	Exercise 1	15
2.1.2	Exercise 2	16
2.1.3	Exercise 3	16
2.1.4	Exercise 4	17
2.1.5	Exercise 5	17
2.1.6	Exercise 6	18
2.2	<code>filter()</code> and its friends	18
2.2.1	Exercise 1	19
2.2.2	Exercise 2	19
2.2.3	Exercise 3	19
2.2.4	Exercise 4	20
2.2.5	Exercise 5	20
2.2.6	Exercise 6	21
2.3	<code>arrange()</code>	21
2.3.1	Exercise 1	21

2.3.2	Exercise 2	22
2.3.3	Exercise 3	22
2.4	<code>mutate()</code> and its friends	23
2.4.1	Exercise 1	23
2.4.2	Exercise 2	24
2.4.3	Exercise 3	24
2.4.4	Exercise 4	25
2.4.5	Exercise 5	25
2.4.6	Exercise 6	26
2.5	<code>summarise()</code>	27
2.5.1	Exercise 1	27
2.5.2	Exercise 2	27
2.5.3	Exercise 3	27
3	Grouping data	29
3.1	<code>group_by()</code>	29
3.1.1	Exercise 1	29
3.1.2	Exercise 2	30
3.1.3	Exercise 3	30
3.1.4	Exercise 4	31
3.1.5	Exercise 5	31
3.1.6	Exercise 6	31
4	Do	33
4.1	<code>do</code>	33
4.1.1	Exercise 1	33
4.1.2	Exercise 2	34
4.1.3	Exercise 3	34
4.1.4	Exercise 4	34
5	Combining data	37
5.1	Joins: <code>inner_join()</code> , <code>left_join()</code> , <code>right_join()</code> , etc.	37
5.1.1	Exercise 1	37
5.1.2	Exercise 2	38

<i>CONTENTS</i>	5
5.1.3 Exercise 3	39
5.1.4 Exercise 4	39
5.1.5 Exercise 5	40
5.1.6 Exercise 6	40
6 Tidy data with tidyr	43
6.1 tidyr	43
6.1.1 Exercise 1	43
6.1.2 Exercise 2	44
6.1.3 Exercise 3	44
6.1.4 Exercise 4	44
7 Handling Missing values	45
7.1 Data import: set working directory	45
7.1.1 Exercise 1	45
7.1.2 Exercise 2	46
7.1.3 Exercise 3	47
8 Dates with lubridate	49
8.1 lubridate	49
8.1.1 Exercise 1	49
8.1.2 Exercise 2	49
8.1.3 Exercise 3	50
9 Manipulating strings with stringr	51
9.1 Data import: set working directory	51
9.1.1 Exercise 1	51
9.1.2 Exercise 2	52
9.1.3 Exercise 3	52
10 Case study	55
10.1 Recap exercise	55
10.2 Exercise	55

Chapter 1

Introduction

In this document you will find some exercises with the `tidyverse` R packages. They are mainly based on the `nycflights13` data, taken from the `nycflights13` package.

1.1 Introduction to `nycflights13` data

The `nycflights13` package contains information about all flights that departed from NYC (e.g. EWR, JFK and LGA) in 2013: 336,776 flights in total.

```
require(nycflights13)
ls(pos = "package:nycflights13")

## [1] "airlines" "airports" "flights"  "planes"   "weather"
```

To help understand what causes delays, it includes a number of useful datasets:

- `flights`: information about all flights that departed from NYC
- `weather`: hourly meteorological data for each airport;
- `planes`: construction information about each plane;
- `airports`: airport names and locations;
- `airlines`: translation between two letter carrier codes and names.

1.1.1 `flights`

This dataset contains on-time data for all flights that departed from NYC (i.e. JFK, LGA or EWR) in 2013. The data frame has 16 variables and 336776 observations. The variables are organised as follow:

- Date of departure: `year`, `month`, `day`;
- Departure and arrival times (local tz): `dep_time`, `arr_time`;
- Departure and arrival delays, in minutes: `dep_delay`, `arr_delay` (negative times represent early departures/arrivals);
- Time of departure broken in to hour and minutes: `hour`, `minute`;
- Two letter carrier abbreviation: `carrier`;
- Plane tail number: `tailnum`;
- Flight number: `flight`;
- Origin and destination: `origin`, `dest`;
- Amount of time spent in the air: `air_time`;
- Distance flown: `distance`.

```
dim(flights)
```

```
## [1] 336776      19
```

```
head(flights)
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517             515         2     830
## 2  2013     1     1     533             529         4     850
## 3  2013     1     1     542             540         2     923
## 4  2013     1     1     544             545        -1    1004
## 5  2013     1     1     554             600        -6     812
## 6  2013     1     1     554             558        -4     740
## # ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

```
str(flights)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   336776 obs. of  19 variables:
## $ year      : int  2013 2013 2013 2013 2013 2013 2013 2013 2013 2013 2013 ...
## $ month     : int   1  1  1  1  1  1  1  1  1  1  1 ...
## $ day       : int   1  1  1  1  1  1  1  1  1  1  1 ...
## $ dep_time  : int  517 533 542 544 554 554 555 557 557 558 ...
## $ sched_dep_time: int  515 529 540 545 600 558 600 600 600 600 ...
```



```
## $ dep_delay      : num  2 4 2 -1 -6 -4 -5 -3 -3 -2 ...
## $ arr_time       : int   830 850 923 1004 812 740 913 709 838 753 ...
## $ sched_arr_time: int   819 830 850 1022 837 728 854 723 846 745 ...
## $ arr_delay      : num   11 20 33 -18 -25 12 19 -14 -8 8 ...
## $ carrier        : chr   "UA" "UA" "AA" "B6" ...
## $ flight         : int  1545 1714 1141 725 461 1696 507 5708 79 301 ...
## $ tailnum        : chr   "N14228" "N24211" "N619AA" "N804JB" ...
## $ origin         : chr   "EWR" "LGA" "JFK" "JFK" ...
## $ dest           : chr   "IAH" "IAH" "MIA" "BQN" ...
## $ air_time       : num   227 227 160 183 116 150 158 53 140 138 ...
## $ distance       : num   1400 1416 1089 1576 762 ...
## $ hour           : num    5 5 5 5 6 5 6 6 6 6 ...
## $ minute         : num   15 29 40 45 0 58 0 0 0 0 ...
## $ time_hour      : POSIXct, format: "2013-01-01 05:00:00" "2013-01-01 05:00:00" ...
```

1.1.2 airlines

This dataset contains airlines names and their respective carrier codes, it has 2 variables and 16 observations. Data structure shows that both variables involved are categorical.

```
dim(airlines)
```

```
## [1] 16  2
```

```
head(airlines)
```

```
## # A tibble: 6 x 2
##   carrier      name
##   <chr>      <chr>
## 1     9E Endeavor Air Inc.
## 2     AA American Airlines Inc.
## 3     AS  Alaska Airlines Inc.
## 4     B6   JetBlue Airways
## 5     DL   Delta Air Lines Inc.
## 6     EV ExpressJet Airlines Inc.
```

```
str(airlines)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   16 obs. of  2 variables:
## $ carrier: chr  "9E" "AA" "AS" "B6" ...
## $ name : chr  "Endeavor Air Inc." "American Airlines Inc." "Alaska Airlines Inc." "JetBlue Airways"
```

1.1.3 airports

This dataset contains useful metadata about airports, that is:

- FAA airport code: `faa`;
- Usual name of the airport: `name`;
- Location of airport: `lat`, `lon`;
- Altitude (in feet): `alt`;
- Timezone offset from GMT: `tz`;
- Daylight savings time zone: `dst` A = Standard US DST: starts on the second Sunday of March, ends on the first Sunday of November U = unknown N = no dst

The data frame has 7 variables and 1397 observations.

```
dim(airports)
```

```
## [1] 1458    8
```

```
head(airports)
```

```
## # A tibble: 6 x 8
##   faa          name      lat      lon    alt    tz
##   <chr>          <chr>    <dbl>    <dbl> <int> <dbl>
## 1   04G      Lansdowne Airport 41.13047 -80.61958 1044   -5
## 2   06A Moton Field Municipal Airport 32.46057 -85.68003 264    -6
## 3   06C      Schaumburg Regional 41.98934 -88.10124 801    -6
## 4   06N      Randall Airport 41.43191 -74.39156 523    -5
## 5   09J      Jekyll Island Airport 31.07447 -81.42778 11     -5
## 6   0A9 Elizabethton Municipal Airport 36.37122 -82.17342 1593   -5
## # ... with 2 more variables: dst <chr>, tzone <chr>
```

```
str(airports)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':    1458 obs. of  8 variables:
## $ faa : chr  "04G" "06A" "06C" "06N" ...
## $ name : chr  "Lansdowne Airport" "Moton Field Municipal Airport" "Schaumburg Regional" "Rand
## $ lat : num  41.1 32.5 42 41.4 31.1 ...
## $ lon : num  -80.6 -85.7 -88.1 -74.4 -81.4 ...
## $ alt : int  1044 264 801 523 11 1593 730 492 1000 108 ...
## $ tz : num  -5 -6 -6 -5 -5 -5 -5 -5 -5 -8 ...
## $ dst : chr  "A" "A" "A" "A" ...
## $ tzone: chr  "America/New_York" "America/Chicago" "America/Chicago" "America/New_York" ...
```

```
## - attr(*, "spec")=List of 2
## ..$ cols :List of 12
## .. ..$ id : list()
## .. ..- attr(*, "class")= chr "collector_integer" "collector"
## .. ..$ name : list()
## .. ..- attr(*, "class")= chr "collector_character" "collector"
## .. ..$ city : list()
## .. ..- attr(*, "class")= chr "collector_character" "collector"
## .. ..$ country: list()
## .. ..- attr(*, "class")= chr "collector_character" "collector"
## .. ..$ faa : list()
## .. ..- attr(*, "class")= chr "collector_character" "collector"
## .. ..$ icao : list()
## .. ..- attr(*, "class")= chr "collector_character" "collector"
## .. ..$ lat : list()
## .. ..- attr(*, "class")= chr "collector_double" "collector"
## .. ..$ lon : list()
## .. ..- attr(*, "class")= chr "collector_double" "collector"
## .. ..$ alt : list()
## .. ..- attr(*, "class")= chr "collector_integer" "collector"
## .. ..$ tz : list()
## .. ..- attr(*, "class")= chr "collector_double" "collector"
## .. ..$ dst : list()
## .. ..- attr(*, "class")= chr "collector_character" "collector"
## .. ..$ tzone : list()
## .. ..- attr(*, "class")= chr "collector_character" "collector"
## ..$ default: list()
## .. ..- attr(*, "class")= chr "collector_guess" "collector"
## ..- attr(*, "class")= chr "col_spec"
```

1.1.4 planes

This dataset contains plane metadata for all plane tailnumbers found in the FAA aircraft registry (American Airways (AA) and Envoy Air (MQ) report fleet numbers rather than tail numbers). The data frame has 9 variables and 3322 observations. The variables are organised as follow:

- Tail number: `tailnum`;
- Year manufactured: `year`;
- Type of plane: `type`;
- Manufacturer and model: `manufacturer`, `model`;
- Number of engines and seats: `engines`, `seats`;
- Average cruising speed in mph: `speed`;
- Type of engine: `engine`.

```

dim(planes)

## [1] 3322    9

head(planes)

## # A tibble: 6 x 9
##   tailnum year      type      manufacturer      model engines
##   <chr> <int>      <chr>      <chr>      <chr> <int>
## 1 N10156  2004 Fixed wing multi engine      EMBRAER EMB-145XR      2
## 2 N102UW  1998 Fixed wing multi engine AIRBUS  INDUSTRIE  A320-214      2
## 3 N103US  1999 Fixed wing multi engine AIRBUS  INDUSTRIE  A320-214      2
## 4 N104UW  1999 Fixed wing multi engine AIRBUS  INDUSTRIE  A320-214      2
## 5 N10575  2002 Fixed wing multi engine      EMBRAER EMB-145LR      2
## 6 N105UW  1999 Fixed wing multi engine AIRBUS  INDUSTRIE  A320-214      2
## # ... with 3 more variables: seats <int>, speed <int>, engine <chr>

str(planes)

## Classes 'tbl_df', 'tbl' and 'data.frame':   3322 obs. of  9 variables:
##  $ tailnum      : chr  "N10156" "N102UW" "N103US" "N104UW" ...
##  $ year         : int   2004 1998 1999 1999 2002 1999 1999 1999 1999 1999 ...
##  $ type         : chr  "Fixed wing multi engine" "Fixed wing multi engine" "Fixed wing multi engine" ...
##  $ manufacturer: chr  "EMBRAER" "AIRBUS INDUSTRIE" "AIRBUS INDUSTRIE" "AIRBUS INDUSTRIE" ...
##  $ model        : chr  "EMB-145XR" "A320-214" "A320-214" "A320-214" ...
##  $ engines       : int   2 2 2 2 2 2 2 2 2 2 ...
##  $ seats        : int  55 182 182 182 55 182 182 182 182 182 ...
##  $ speed        : int   NA NA NA NA NA NA NA NA NA NA ...
##  $ engine       : chr  "Turbo-fan" "Turbo-fan" "Turbo-fan" "Turbo-fan" ...

```

1.1.5 weather

This dataset is about hourly meteorological data for LGA, JFK and EWR. The data frame has 14 variables and 8719 observations. The variables are organised as follow:

- Weather station: `origin` (named `origin` to facilitate merging with `flights` data);
- Time of recording: `year`, `month`, `day`, `hour`;
- Temperature and dewpoint in F: `temp`, `dewp`;
- Relative humidity: `humid`;
- Wind direction (in degrees), speed and gust speed (in mph): `wind_dir`, `wind_speed`, `wind_gust`;

- Precipitation, in inches: `precip`;
- Sea level pressure in millibars: `pressure`;
- Visibility in miles: `visib`.

```
dim(weather)
```

```
## [1] 26130    15
```

```
head(weather)
```

```
## # A tibble: 6 x 15
##   origin year month   day hour  temp  dewp humid wind_dir wind_speed
##   <chr> <dbl> <dbl> <int> <int> <dbl> <dbl> <dbl>   <dbl>     <dbl>
## 1   EWR  2013     1     1     0 37.04 21.92 53.97     230    10.35702
## 2   EWR  2013     1     1     1 37.04 21.92 53.97     230    13.80936
## 3   EWR  2013     1     1     2 37.94 21.92 52.09     230    12.65858
## 4   EWR  2013     1     1     3 37.94 23.00 54.51     230    13.80936
## 5   EWR  2013     1     1     4 37.94 24.08 57.04     240    14.96014
## 6   EWR  2013     1     1     6 39.02 26.06 59.37     270    10.35702
## # ... with 5 more variables: wind_gust <dbl>, precip <dbl>,
## #   pressure <dbl>, visib <dbl>, time_hour <dtm>
```

```
str(weather)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   26130 obs. of  15 variables:
## $ origin   : chr  "EWR" "EWR" "EWR" "EWR" ...
## $ year     : num  2013 2013 2013 2013 2013 ...
## $ month    : num  1 1 1 1 1 1 1 1 1 1 ...
## $ day      : int  1 1 1 1 1 1 1 1 1 1 ...
## $ hour     : int  0 1 2 3 4 6 7 8 9 10 ...
## $ temp     : num  37 37 37.9 37.9 37.9 ...
## $ dewp     : num  21.9 21.9 21.9 23 24.1 ...
## $ humid    : num  54 54 52.1 54.5 57 ...
## $ wind_dir : num  230 230 230 230 240 270 250 240 250 260 ...
## $ wind_speed: num  10.4 13.8 12.7 13.8 15 ...
## $ wind_gust : num  11.9 15.9 14.6 15.9 17.2 ...
## $ precip   : num  0 0 0 0 0 0 0 0 0 0 ...
## $ pressure : num  1014 1013 1013 1013 1013 ...
## $ visib    : num  10 10 10 10 10 10 10 10 10 10 ...
## $ time_hour : POSIXct, format: "2013-01-01 01:00:00" "2013-01-01 02:00:00" ...
```


Chapter 2

Verb functions

In this section you will find exercises on the basic verbs of data manipulating provided by `dplyr`:

1. `select()`;
2. `filter()`;
3. `arrange()`;
4. `mutate()`;
5. `summarise()`.

2.1 `select()` and its friends

Note: all the exercises of this section are based on the `flights` dataset.

```
require(tidyverse)
require(nycflights13)
```

2.1.1 Exercise 1

Extract the following information about flights:

- `month`;
- `day`;
- `air_time`;
- `distance`.

```
# require(nycflights13)
data(flights)

flights %>% select(month, day, air_time, distance)

## # A tibble: 336,776 x 4
##   month   day air_time distance
##   <int> <int>   <dbl>   <dbl>
## 1     1     1     227     1400
## 2     1     1     227     1416
## 3     1     1     160     1089
## 4     1     1     183     1576
## 5     1     1     116      762
## 6     1     1     150      719
## 7     1     1     158     1065
## 8     1     1      53      229
## 9     1     1     140      944
## 10    1     1     138      733
## # ... with 336,766 more rows
```

2.1.2 Exercise 2

Extract all information about flights except hour and minute.

```
flights %>% select(-hour, -minute)

## # A tibble: 336,776 x 17
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830             819
## 2  2013     1     1     533             529           4     850             830
## 3  2013     1     1     542             540           2     923             850
## 4  2013     1     1     544             545          -1    1004            1022
## 5  2013     1     1     554             600          -6     812             837
## 6  2013     1     1     554             558          -4     740             728
## 7  2013     1     1     555             600          -5     913             854
## 8  2013     1     1     557             600          -3     709             723
## 9  2013     1     1     557             600          -3     838             846
## 10 2013     1     1     558             600          -2     753             745
## # ... with 336,766 more rows, and 9 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, time_hour <dtm>
```

2.1.3 Exercise 3

Select all variables whose name ends in “time”.


```
flights %>% select(ends_with("time"))
```

```
## # A tibble: 336,776 x 5
```

	dep_time	sched_dep_time	arr_time	sched_arr_time	air_time
	<int>	<int>	<int>	<int>	<dbl>
## 1	517	515	830	819	227
## 2	533	529	850	830	227
## 3	542	540	923	850	160
## 4	544	545	1004	1022	183
## 5	554	600	812	837	116
## 6	554	558	740	728	150
## 7	555	600	913	854	158
## 8	557	600	709	723	53
## 9	557	600	838	846	140
## 10	558	600	753	745	138

```
## # ... with 336,766 more rows
```

2.1.4 Exercise 4

Select all variables whose name contains the word “delay”.

```
flights %>% select(matches("delay"))
```

```
## # A tibble: 336,776 x 2
```

	dep_delay	arr_delay
	<dbl>	<dbl>
## 1	2	11
## 2	4	20
## 3	2	33
## 4	-1	-18
## 5	-6	-25
## 6	-4	12
## 7	-5	19
## 8	-3	-14
## 9	-3	-8
## 10	-2	8

```
## # ... with 336,766 more rows
```

2.1.5 Exercise 5

Select the `tailnum` variable and rename it into `tail_num`.

```
flights %>% select(tail_num = tailnum)
```

```
## # A tibble: 336,776 x 1
##   tail_num
##   <chr>
## 1 N14228
## 2 N24211
## 3 N619AA
## 4 N804JB
## 5 N668DN
## 6 N39463
## 7 N516JB
## 8 N829AS
## 9 N593JB
## 10 N3ALAA
## # ... with 336,766 more rows
```

2.1.6 Exercise 6

Select all the variables and rename the `tailnum` variable into `tail_num`.

```
flights %>% rename(tail_num = tailnum)

## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830           819
## 2  2013     1     1     533             529           4     850           830
## 3  2013     1     1     542             540           2     923           850
## 4  2013     1     1     544             545          -1    1004          1022
## 5  2013     1     1     554             600          -6     812           837
## 6  2013     1     1     554             558          -4     740           728
## 7  2013     1     1     555             600          -5     913           854
## 8  2013     1     1     557             600          -3     709           723
## 9  2013     1     1     557             600          -3     838           846
## 10 2013     1     1     558             600          -2     753           745
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tail_num <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

2.2 filter() and its friends

Note: all the exercises of this section are based on the `flights` dataset.

```
require(tidyverse)
require(nycflights13)
```

2.2.1 Exercise 1

Select all flights which delayed more than 1000 minutes at departure.

```
flights %>% filter(dep_delay > 1000)

## # A tibble: 5 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     9     641             900         1301    1242         1530
## 2  2013     1    10    1121            1635         1126    1239         1810
## 3  2013     6    15    1432            1935         1137    1607         2120
## 4  2013     7    22     845            1600         1005    1044         1815
## 5  2013     9    20    1139            1845         1014    1457         2210
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

2.2.2 Exercise 2

Select all flights which delayed more than 1000 minutes at departure or at arrival.

```
flights %>% filter(dep_delay > 1000 | arr_delay > 1000)

## # A tibble: 5 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     9     641             900         1301    1242         1530
## 2  2013     1    10    1121            1635         1126    1239         1810
## 3  2013     6    15    1432            1935         1137    1607         2120
## 4  2013     7    22     845            1600         1005    1044         1815
## 5  2013     9    20    1139            1845         1014    1457         2210
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# alternatively
# flights %>% filter(dep_delay > 1000, arr_delay > 1000)
```

2.2.3 Exercise 3

Select all flights which took off from “EWR” and landed in “IAH” on Christmas Day.

```
flights %>% filter(origin == "EWR" & dest == "IAH" & month == 12 & day == 25)
```

```
## # A tibble: 8 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013    12    25     524           515         9      805           814
## 2  2013    12    25     753           747         6     1038          1048
## 3  2013    12    25    1018          1015         3     1310          1316
## 4  2013    12    25    1442          1345        57     1730          1646
## 5  2013    12    25    1530          1529         1     1836          1826
## 6  2013    12    25    1628          1630        -2     1944          1925
## 7  2013    12    25    1843          1804        39     2141          2113
## 8  2013    12    25    2003          2006        -3     2304          2314
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>

# alternatively
# flights %>% filter(origin == "EWR", dest == "IAH", month == 12, day ==25)
```

2.2.4 Exercise 4

Select the first five flights in this dataset.

```
flights %>% slice(1:5)

## # A tibble: 5 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2      830           819
## 2  2013     1     1     533           529         4      850           830
## 3  2013     1     1     542           540         2      923           850
## 4  2013     1     1     544           545        -1     1004          1022
## 5  2013     1     1     554           600        -6      812           837
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

2.2.5 Exercise 5

Select the last ten flights in this dataset.

```
flights %>% slice((n()-9):n())

## # A tibble: 10 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
```

```
##      <int> <int> <int>      <int>      <int>      <dbl>      <int>      <int>
## 1  2013      9    30      2240      2250      -10      2347          7
## 2  2013      9    30      2241      2246       -5      2345          1
## 3  2013      9    30      2307      2255       12      2359      2358
## 4  2013      9    30      2349      2359      -10       325      350
## 5  2013      9    30         NA      1842        NA        NA      2019
## 6  2013      9    30         NA      1455        NA        NA      1634
## 7  2013      9    30         NA      2200        NA        NA      2312
## 8  2013      9    30         NA      1210        NA        NA      1330
## 9  2013      9    30         NA      1159        NA        NA      1344
## 10 2013      9    30         NA       840        NA        NA      1020
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

2.2.6 Exercise 6

Extract information about distance for all flights which delayed more than 1000 minutes at departure.

```
flights %>%
  filter(dep_delay > 1000) %>%
  select(distance)

## # A tibble: 5 x 1
##   distance
##   <dbl>
## 1    4983
## 2     719
## 3     483
## 4     589
## 5    2586
```

2.3 arrange()

Note: all the exercises of this section are based on the `flights` dataset.

```
require(tidyverse)
require(nycflights13)
```

2.3.1 Exercise 1

Sort the flights in chronological order.

```
flights %>% arrange(year, month, day)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## 9  2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

2.3.2 Exercise 2

Sort the flights by decreasing arrival delay.

```
flights %>% arrange(desc(arr_delay))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     9     641           900        1301    1242          1530
## 2  2013     6    15    1432          1935        1137    1607          2120
## 3  2013     1    10    1121          1635        1126    1239          1810
## 4  2013     9    20    1139          1845        1014    1457          2210
## 5  2013     7    22     845          1600        1005    1044          1815
## 6  2013     4    10    1100          1900         960    1342          2211
## 7  2013     3    17    2321           810         911     135          1020
## 8  2013     7    22    2257           759         898     121          1026
## 9  2013    12     5     756          1700         896    1058          2020
## 10 2013     5     3    1133          2055         878    1250          2215
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

2.3.3 Exercise 3

Sort the flights by origin (in alphabetical order) and decreasing arrival delay.

```
flights %>% arrange(origin, desc(arr_delay))

## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1    10    1121           1635     1126    1239           1810
## 2  2013    12     5     756           1700      896    1058           2020
## 3  2013     5     3    1133           2055      878    1250           2215
## 4  2013    12    19     734           1725      849    1046           2039
## 5  2013    12    17     705           1700      845    1026           2020
## 6  2013    11     3     603           1645      798     829           1913
## 7  2013     2    24    1921           615       786    2135           842
## 8  2013    10    14    2042           900       702    2255          1127
## 9  2013     7    21    1555           615       580    1955           910
##10  2013     7     7    2123          1030      653     17           1345
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

2.4 mutate() and its friends

Note: all the exercises of this section are based on the `flights` dataset.
Times are in minutes and distances are in miles.

```
require(tidyverse)
require(nycflights13)
```

2.4.1 Exercise 1

Add the following new variables to the `flights` dataset:

- the gained time in minutes, defined as the difference between delay at departure and delay at arrival;
- the speed in miles per hour (`distance / air_time * 60`).

Show only the following variables: delay at departure, delay at arrival, distance, air time and the two new variables (gained time and speed).

```
flights %>% mutate(gained_time = arr_delay - dep_delay, speed = distance/air_time*60)

## # A tibble: 336,776 x 21
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
```

```
## 1 2013 1 1 517 515 2 830 819
## 2 2013 1 1 533 529 4 850 830
## 3 2013 1 1 542 540 2 923 850
## 4 2013 1 1 544 545 -1 1004 1022
## 5 2013 1 1 554 600 -6 812 837
## 6 2013 1 1 554 558 -4 740 728
## 7 2013 1 1 555 600 -5 913 854
## 8 2013 1 1 557 600 -3 709 723
## 9 2013 1 1 557 600 -3 838 846
## 10 2013 1 1 558 600 -2 753 745
## # ... with 336,766 more rows, and 13 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,
## #   gained_time <dbl>, speed <dbl>
```

2.4.2 Exercise 2

Redo the previous calculations keeping only the new variables.

```
flights %>%
  transmute(gained_time = arr_delay - dep_delay, speed = distance/air_time*60)

## # A tibble: 336,776 x 2
##   gained_time speed
##   <dbl> <dbl>
## 1 9 370.0441
## 2 16 374.2731
## 3 31 408.3750
## 4 -17 516.7213
## 5 -19 394.1379
## 6 16 287.6000
## 7 24 404.4304
## 8 -11 259.2453
## 9 -5 404.5714
## 10 10 318.6957
## # ... with 336,766 more rows
```

2.4.3 Exercise 3

After sorting flights in chronological order, for each flight calculate the difference between its delay at arrival and the delay at arrival of the immediately previous flight. Have R to show only the delay variables (delay at departure, delay at arrival and the new variable).

```
flights %>%
  arrange(year, month, day) %>%
```



```
mutate(lead_arr_delay = lead(arr_delay), delta_delay = lead_arr_delay - arr_delay) %>%
select(dep_delay, arr_delay, delta_delay)

## # A tibble: 336,776 x 3
##   dep_delay arr_delay delta_delay
##   <dbl>     <dbl>     <dbl>
## 1         2         11          9
## 2         4         20         13
## 3         2         33        -51
## 4        -1        -18         -7
## 5        -6        -25         37
## 6        -4         12          7
## 7        -5         19        -33
## 8        -3        -14          6
## 9        -3         -8         16
## 10       -2          8        -10
## # ... with 336,766 more rows
```

2.4.4 Exercise 4

For each flight calculate the ‘min ranking’ in terms of delay at arrival.

```
flights %>%
  mutate(min_rank_arr_delay = min_rank(arr_delay))

## # A tibble: 336,776 x 20
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515         2     830             819
## 2  2013     1     1     533             529         4     850             830
## 3  2013     1     1     542             540         2     923             850
## 4  2013     1     1     544             545        -1    1004            1022
## 5  2013     1     1     554             600        -6     812             837
## 6  2013     1     1     554             558        -4     740             728
## 7  2013     1     1     555             600        -5     913             854
## 8  2013     1     1     557             600        -3     709             723
## 9  2013     1     1     557             600        -3     838             846
## 10 2013     1     1     558             600        -2     753             745
## # ... with 336,766 more rows, and 12 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,
## #   min_rank_arr_delay <int>
```

2.4.5 Exercise 5

For each flight calculate the ‘first ranking’ in terms of delay at arrival.

```

flights %>%
  mutate(first_rank_arr_delay = row_number(arr_delay))

## # A tibble: 336,776 x 20
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## 9  2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # ... with 336,766 more rows, and 12 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,
## #   first_rank_arr_delay <int>

```

2.4.6 Exercise 6

Create a variable which indicates if a flight took off on time, i.e. departure delay is more than -4 and less than 4 minutes late.

```

flights %>%
  filter(arr_delay > -4 & arr_delay < 4) %>%
  mutate(dep_on_time = 1)

## # A tibble: 37,061 x 20
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     558           600          -2     849           851
## 2  2013     1     1     558           600          -2     853           856
## 3  2013     1     1     622           630          -8    1017          1014
## 4  2013     1     1     623           627          -4     933           932
## 5  2013     1     1     627           630          -3    1018          1018
## 6  2013     1     1     628           630          -2    1137          1140
## 7  2013     1     1     658           700          -2    1027          1025
## 8  2013     1     1     659           700          -1    1008          1007
## 9  2013     1     1     728           732          -4    1041          1038
## 10 2013     1     1     732           735          -3     857           858
## # ... with 37,051 more rows, and 12 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,
## #   dep_on_time <dbl>

```

2.5 summarise()

Note: all the exercises of this section are based on the `flights` dataset.

```
require(tidyverse)
require(nycflights13)
```

2.5.1 Exercise 1

Calculate minimum, mean and maximum delay at arrival.

```
flights %>% summarise(min_arr_delay = min(arr_delay, na.rm = T),
                      mean_arr_delay = mean(arr_delay, na.rm = T),
                      max_arr_delay = max(arr_delay, na.rm = T))

## # A tibble: 1 x 3
##   min_arr_delay mean_arr_delay max_arr_delay
##           <dbl>           <dbl>         <dbl>
## 1          -86           6.895377         1272
```

2.5.2 Exercise 2

Calculate minimum, mean and maximum delay at arrival for flights in January.

```
flights %>%
  filter(month == 1) %>%
  summarise(min_arr_delay = min(arr_delay, na.rm = T),
            mean_arr_delay = mean(arr_delay, na.rm = T),
            max_arr_delay = max(arr_delay, na.rm = T))

## # A tibble: 1 x 3
##   min_arr_delay mean_arr_delay max_arr_delay
##           <dbl>           <dbl>         <dbl>
## 1          -70           6.129972         1272
```

2.5.3 Exercise 3

Calculate the number of flights are contained in the dataset

```
flights %>%
  summarise(n_flights = n())
```

```
## # A tibble: 1 x 1
##   n_flights
##       <int>
## 1     336776
```

Chapter 3

Grouping data

3.1 group_by()

Note: all the exercises of this section are based on the `flights` dataset.

```
library(tidyverse)
library(nycflights13)
```

3.1.1 Exercise 1

Calculate number of flights, minimum, mean and maximum delay at arrival for flights by month.

```
flights %>%
  group_by(month) %>%
  summarise(n_flights = n(),
            min_arr_delay = min(arr_delay, na.rm = T),
            mean_arr_delay = mean(arr_delay, na.rm = T),
            max_arr_delay = max(arr_delay, na.rm = T))
```



```
## # A tibble: 12 x 5
##   month n_flights min_arr_delay mean_arr_delay max_arr_delay
##   <int>   <int>         <dbl>         <dbl>         <dbl>
## 1     1     27004          -70         6.1299720       1272
## 2     2     24951          -70         5.6130194        834
## 3     3     28834          -68         5.8075765        915
## 4     4     28330          -68        11.1760630        931
## 5     5     28796          -86         3.5215088        875
## 6     6     28243          -64        16.4813296       1127
## 7     7     29425          -66        16.7113067        989
## 8     8     29327          -68         6.0406524        490
## 9     9     27574          -68        -4.0183636       1007
## 10    10     28889          -61        -0.1670627        688
```

```
## 11    11    27268          -67      0.4613474      796
## 12    12    28135          -68     14.8703553      878
```

3.1.2 Exercise 2

Calculate number of flights, mean delay at departure and arrival for flights by origin.

```
flights %>%
  group_by(origin) %>%
  summarise(n_flights = n(),
            min_arr_delay = min(arr_delay, na.rm = T),
            mean_arr_delay = mean(arr_delay, na.rm = T),
            max_arr_delay = max(arr_delay, na.rm = T))

## # A tibble: 3 x 5
##   origin n_flights min_arr_delay mean_arr_delay max_arr_delay
##   <chr>   <int>         <dbl>         <dbl>         <dbl>
## 1   EWR    120835         -86          9.107055         1109
## 2   JFK    111279         -79          5.551481         1272
## 3   LGA    104662         -68          5.783488          915
```

3.1.3 Exercise 3

Calculate the number of flights that go to each possible destination.

```
flights %>%
  group_by(dest) %>%
  summarise(n_flights = n())

## # A tibble: 105 x 2
##   dest n_flights
##   <chr>   <int>
## 1   ABQ       254
## 2   ACK       265
## 3   ALB       439
## 4   ANC         8
## 5   ATL    17215
## 6   AUS     2439
## 7   AVL       275
## 8   BDL       443
## 9   BGR       375
## 10  BHM       297
## # ... with 95 more rows
```

3.1.4 Exercise 4

Calculate the number of flights for each day. Save the result in a data frame called `per_day`.

```
per_day <- flights %>%  
  group_by(year, day, month) %>%  
  summarise(n_flights = n())
```

3.1.5 Exercise 5

By exploiting `per_day`, calculate the number of flights for each month. Save the result in a data frame called `per_month`.

```
per_month <- flights %>%  
  group_by(month) %>%  
  summarise(n_flights = n())
```

3.1.6 Exercise 6

Calculate the mean daily number of flights per month.

```
per_month %>%  
  group_by(month) %>%  
  summarise(mean_n_flights = mean(n_flights))
```

```
## # A tibble: 12 x 2  
##   month mean_n_flights  
##   <int>         <dbl>  
## 1     1         27004  
## 2     2         24951  
## 3     3         28834  
## 4     4         28330  
## 5     5         28796  
## 6     6         28243  
## 7     7         29425  
## 8     8         29327  
## 9     9         27574  
## 10    10         28889  
## 11    11         27268  
## 12    12         28135
```


Chapter 4

Do

4.1 do

Note: all the exercises of this section are based on the `flights` dataset.

```
library(dplyr)
library(nycflights13)
```

4.1.1 Exercise 1

Calculate quartiles (25-, 50- and 75-percentiles) of delay at arrival per origin. Put all three quartiles in a unique column.

```
flights %>% group_by(origin) %>%
  do(data.frame(p = (1:3)/4,
               quantile = quantile(.$arr_delay, probs = (1:3)/4, na.rm = TRUE)))

## # A tibble: 9 x 3
## # Groups:   origin [3]
##   origin      p quantile
##   <chr> <dbl>   <dbl>
## 1   EWR 0.25    -16
## 2   EWR 0.50     -4
## 3   EWR 0.75     16
## 4   JFK 0.25    -18
## 5   JFK 0.50     -6
## 6   JFK 0.75     13
## 7   LGA 0.25    -17
## 8   LGA 0.50     -5
## 9   LGA 0.75     12
```

4.1.2 Exercise 2

Redo the previous exercise putting the three quartiles in three different columns (hint: use `summarise()`).

```
flights %>% group_by(origin) %>%
  do(data.frame(p1 = quantile(.$arr_delay, probs = 1/4, na.rm = TRUE),
    p2 = quantile(.$arr_delay, probs = 2/4, na.rm = TRUE),
    p3 = quantile(.$arr_delay, probs = 3/4, na.rm = TRUE)))

## # A tibble: 3 x 4
## # Groups:   origin [3]
##   origin    p1    p2    p3
##   <chr> <dbl> <dbl> <dbl>
## 1   EWR   -16    -4   16
## 2   JFK   -18    -6   13
## 3   LGA   -17    -5   12
```

4.1.3 Exercise 3

Calculate mean and standard deviation of delay at arrival per origin. Put both statistics in a unique column.

```
fun <- function(x, ...) c(mean = mean(x, ...), sd = sd(x, ...))

flights %>% group_by(origin) %>%
  do(data.frame(stats = c("mean", "sd") ,
    value = fun(.$arr_delay, na.rm = TRUE)))

## # A tibble: 6 x 3
## # Groups:   origin [3]
##   origin stats      value
##   <chr> <fctr>      <dbl>
## 1   EWR   mean  9.107055
## 2   EWR    sd 45.529183
## 3   JFK   mean  5.551481
## 4   JFK    sd 44.277448
## 5   LGA   mean  5.783488
## 6   LGA    sd 43.862273
```

4.1.4 Exercise 4

Redo the previous exercise putting mean and standard deviation in two different columns (hint: use `summarise()`).

```
flights %>% group_by(origin) %>%  
  do(data.frame(arr_delay_mean = mean(.$arr_delay, na.rm = TRUE),  
                arr_delay_sd = sd(.$arr_delay, na.rm = TRUE)))  
  
## # A tibble: 3 x 3  
## # Groups:   origin [3]  
##   origin arr_delay_mean arr_delay_sd  
##   <chr>         <dbl>         <dbl>  
## 1 EWR          9.107055      45.52918  
## 2 JFK          5.551481      44.27745  
## 3 LGA          5.783488      43.86227
```


Chapter 5

Combining data

5.1 Joins: `inner_join()`, `left_join()`, `right_join()`, etc.

Note: all the exercises of this section are based on `flights`, `airlines`, `airports` or `planes` datasets.

```
library(dplyr)
library(nycflights13)
```

5.1.1 Exercise 1

Keep only the following variables of the `flights` dataset: `month`, `day`, `hour`, `origin`, `destination` and `carrier`. Save this dataset in a data frame and call it `flights_red`. Through a proper join command, add the carrier name to `flights_red` (this piece of information is available in `airlines`).

```
flights_red <- flights %>%
  select(month, day, hour, origin, dest, carrier)

right_join(flights_red, airlines)

## Joining, by = "carrier"

## # A tibble: 336,776 x 7
##   month   day hour origin dest carrier      name
##   <int> <int> <dbl> <chr> <chr> <chr>    <chr>
## 1     1     1     8   JFK   MSP    9E Endeavor Air Inc.
## 2     1     1    15   JFK   IAD    9E Endeavor Air Inc.
## 3     1     1    14   JFK   BUF    9E Endeavor Air Inc.
## 4     1     1    15   JFK   SYR    9E Endeavor Air Inc.
```

```
## 5      1      1     15   JFK   ROC      9E Endeavor Air Inc.
## 6      1      1     15   JFK   BWI      9E Endeavor Air Inc.
## 7      1      1     15   JFK   ORD      9E Endeavor Air Inc.
## 8      1      1     15   JFK   IND      9E Endeavor Air Inc.
## 9      1      1     16   JFK   BNA      9E Endeavor Air Inc.
## 10     1      1     16   JFK   BOS      9E Endeavor Air Inc.
## # ... with 336,766 more rows
```

5.1.2 Exercise 2

Through a proper join command, add name, latitude, longitude and altitude of the origin airport to `flights_red` (these pieces of information are available in `airports`). Do the same also for the destination airport. (If you are able to, try to keep variables about both origin and destination airports in the same final dataset).

```
flights_red %>% left_join(airports, c("origin" = "faa"))
```

```
## # A tibble: 336,776 x 13
##   month   day   hour origin dest carrier      name      lat
##   <int> <int> <dbl> <chr> <chr> <chr>      <chr>    <dbl>
## 1     1     1     5   EWR   IAH    UA Newark Liberty Intl 40.69250
## 2     1     1     5   LGA   IAH    UA      La Guardia 40.77725
## 3     1     1     5   JFK   MIA    AA John F Kennedy Intl 40.63975
## 4     1     1     5   JFK   BQN    B6 John F Kennedy Intl 40.63975
## 5     1     1     6   LGA   ATL    DL      La Guardia 40.77725
## 6     1     1     5   EWR   ORD    UA Newark Liberty Intl 40.69250
## 7     1     1     6   EWR   FLL    B6 Newark Liberty Intl 40.69250
## 8     1     1     6   LGA   IAD    EV      La Guardia 40.77725
## 9     1     1     6   JFK   MCO    B6 John F Kennedy Intl 40.63975
## 10    1     1     6   LGA   ORD    AA      La Guardia 40.77725
## # ... with 336,766 more rows, and 5 more variables: lon <dbl>, alt <int>,
## #   tz <dbl>, dst <chr>, tzone <chr>
```

```
flights_red %>% left_join(airports, c("dest" = "faa"))
```

```
## # A tibble: 336,776 x 13
##   month   day   hour origin dest carrier      name
##   <int> <int> <dbl> <chr> <chr> <chr>      <chr>
## 1     1     1     5   EWR   IAH    UA George Bush Intercontinental
## 2     1     1     5   LGA   IAH    UA George Bush Intercontinental
## 3     1     1     5   JFK   MIA    AA Miami Intl
## 4     1     1     5   JFK   BQN    B6 <NA>
## 5     1     1     6   LGA   ATL    DL Hartsfield Jackson Atlanta Intl
## 6     1     1     5   EWR   ORD    UA Chicago Ohare Intl
## 7     1     1     6   EWR   FLL    B6 Fort Lauderdale Hollywood Intl
```

```
## 8      1      1      6    LGA    IAD      EV      Washington Dulles Intl
## 9      1      1      6    JFK    MCO      B6              Orlando Intl
## 10     1      1      6    LGA    ORD      AA              Chicago Ohare Intl
## # ... with 336,766 more rows, and 6 more variables: lat <dbl>, lon <dbl>,
## #   alt <int>, tz <dbl>, dst <chr>, tzone <chr>
```

5.1.3 Exercise 3

Through the `inner_join()` function, redo the same for the destination airport but keep only the flights whose information is available in both datasets (`flights` and `airports`).

```
flights_red %>% inner_join(airports, c("dest" = "faa"))

## # A tibble: 329,174 x 13
##   month   day hour origin dest carrier      name
##   <int> <int> <dbl> <chr> <chr>   <chr>   <chr>
## 1     1     1     5    EWR   IAH     UA   George Bush Intercontinental
## 2     1     1     5    LGA   IAH     UA   George Bush Intercontinental
## 3     1     1     5    JFK   MIA     AA           Miami Intl
## 4     1     1     6    LGA   ATL     DL Hartsfield Jackson Atlanta Intl
## 5     1     1     5    EWR   ORD     UA           Chicago Ohare Intl
## 6     1     1     6    EWR   FLL     B6   Fort Lauderdale Hollywood Intl
## 7     1     1     6    LGA   IAD     EV      Washington Dulles Intl
## 8     1     1     6    JFK   MCO     B6              Orlando Intl
## 9     1     1     6    LGA   ORD     AA           Chicago Ohare Intl
## 10    1     1     6    JFK   PBI     B6           Palm Beach Intl
## # ... with 329,164 more rows, and 6 more variables: lat <dbl>, lon <dbl>,
## #   alt <int>, tz <dbl>, dst <chr>, tzone <chr>
```

5.1.4 Exercise 4

Redo the exercise 3 by using `full_join()` instead of `inner_join()`. What is the difference in the result?

```
flights_red %>% full_join(airports, c("dest" = "faa"))

## # A tibble: 338,133 x 13
##   month   day hour origin dest carrier      name
##   <int> <int> <dbl> <chr> <chr>   <chr>   <chr>
## 1     1     1     5    EWR   IAH     UA   George Bush Intercontinental
## 2     1     1     5    LGA   IAH     UA   George Bush Intercontinental
## 3     1     1     5    JFK   MIA     AA           Miami Intl
## 4     1     1     5    JFK   BQN     B6              <NA>
## 5     1     1     6    LGA   ATL     DL Hartsfield Jackson Atlanta Intl
```

```
## 6      1      1      5    EWR    ORD      UA      Chicago Ohare Intl
## 7      1      1      6    EWR    FLL      B6    Fort Lauderdale Hollywood Intl
## 8      1      1      6    LGA    IAD      EV      Washington Dulles Intl
## 9      1      1      6    JFK    MCO      B6      Orlando Intl
## 10     1      1      6    LGA    ORD      AA      Chicago Ohare Intl
## # ... with 338,123 more rows, and 6 more variables: lat <dbl>, lon <dbl>,
## #   alt <int>, tz <dbl>, dst <chr>, tzone <chr>

# there are a few more rows due to the fact that full_join keeps all
# rows even those with no matches
```

5.1.5 Exercise 5

Through the `anti_join()` function, extract all the flights from `flights` whose information about destination airport is not available in `airports`.

```
flights_red %>% anti_join(airports, c("dest" = "faa"))

## # A tibble: 7,602 x 6
##   month   day hour origin dest carrier
##   <int> <int> <dbl> <chr> <chr>   <chr>
## 1     1     1     23   JFK   PSE     B6
## 2     1     2     23   JFK   PSE     B6
## 3     1     3     23   JFK   PSE     B6
## 4     1     4     23   JFK   PSE     B6
## 5     1     5     23   JFK   PSE     B6
## 6     1     6     23   JFK   PSE     B6
## 7     1     7     23   JFK   PSE     B6
## 8     1     8     23   JFK   PSE     B6
## 9     1     9     23   JFK   PSE     B6
## 10    1    10     23   JFK   PSE     B6
## # ... with 7,592 more rows
```

5.1.6 Exercise 6

Sort the `planes` dataset by increasing year. Then create two datasets: the first will deal with planes older than 2000; the second will deal with planes of 2000 or newer. Finally create a unique dataset where the first rows will deal with the newest planes, whereas the last rows will deal with the oldest planes.

```
planes_old <- planes %>%
  arrange(year) %>%
  filter(year <= 2000)
```



```

planes_young <- planes %>%
  arrange(year) %>%
  filter(year > 2000)

planes_old %>% bind_rows(planes_young)

## # A tibble: 3,252 x 9
##   tailnum year      type manufacturer      model engines
##   <chr> <int>      <chr>      <chr>      <chr>   <int>
## 1 N381AA 1956 Fixed wing multi engine    DOUGLAS    DC-7BF         4
## 2 N201AA 1959 Fixed wing single engine    CESSNA      150         1
## 3 N567AA 1959 Fixed wing single engine DEHAVILLAND OTTER DHC-3         1
## 4 N378AA 1963 Fixed wing single engine    CESSNA     172E         1
## 5 N575AA 1963 Fixed wing single engine    CESSNA    210-5(205)         1
## 6 N14629 1965 Fixed wing multi engine    BOEING     737-524         2
## 7 N615AA 1967 Fixed wing multi engine    BEECH      65-A90         2
## 8 N425AA 1968 Fixed wing single engine    PIPER     PA-28-180         1
## 9 N383AA 1972 Fixed wing multi engine    BEECH      E-90           2
## 10 N364AA 1973 Fixed wing multi engine    CESSNA     310Q           2
## # ... with 3,242 more rows, and 3 more variables: seats <int>, speed <int>,
## #   engine <chr>

```


Chapter 6

Tidy data with tidyr

6.1 tidyr

```
library(tidyverse)
```

6.1.1 Exercise 1

Consider the following dataset:

```
heartrate_wide <- data.frame(
  name = c("Aldo", "Giovanni", "Giacomo"),
  surname = c("Baglio", "Storti", "Poretti"),
  morning = c(67, 80, 64),
  afternoon = c(56, 90, 50)
)
heartrate_wide
```

	name	surname	morning	afternoon
## 1	Aldo	Baglio	67	56
## 2	Giovanni	Storti	80	90
## 3	Giacomo	Poretti	64	50

It represents the heart rate measured on three patients in the morning and in the afternoon. The dataset is in the wide format: change it to the long format through a proper `tidyr` function. Save the result in a data frame and call it `heartrate_long`.

```
heartrate_long <- gather(heartrate_wide, key = "when", value = "value" , 3:4)
```

6.1.2 Exercise 2

Starting from `heartrate_long`, come back to a dataset in a wide format through a proper `tidyr` function. The result should be obviously equal to `heartrate_wide`.

```
spread(heartrate_long, key = when, value = "value")
```

```
##      name surname afternoon morning
## 1    Aldo  Baglio         56       67
## 2  Giacomo Poretti         50       64
## 3 Giovanni Storti         90       80
```

6.1.3 Exercise 3

Consider the dataset `heartrate_wide` and unite name and surname of the patients in a unique column through a proper `tidyr` function. Save the result in a new data frame called `heartrate_united`.

```
heartrate_united <- heartrate_wide %>%
  unite(name_surname, name, surname)
```

6.1.4 Exercise 4

Starting from `heartrate_united`, come back to a dataset where name and surname are in two different columns through a proper `tidyr` function. The result should be obviously equal to `heartrate_wide`.

```
heartrate_united %>%
  separate(name_surname, c("name", "surname"))
```

```
##      name surname morning afternoon
## 1    Aldo  Baglio         67         56
## 2 Giovanni Storti         80         90
## 3  Giacomo Poretti         64         50
```

Chapter 7

Handling Missing values

7.1 Data import: set working directory

Some of the data that will be used in this exercises are contained in the data folder. Hence you should set your working directory in the *data* folder, using `setwd()` function, like in this example:

```
setwd("~/Documents/datamanager/exercises/data")
```

You will work inside this folder.

```
library(tidyverse)
```

7.1.1 Exercise 1

Consider the following dataset:

```
heartrate <- data.frame(
  name = c("Aldo", "Giovanni", "Giacomo", "Aldo", "Giovanni", "Giacomo",
           "Giovanni", "Giacomo"),
  surname = c("Baglio", "Storti", "Poretti", "Baglio", "Storti", "Poretti",
              "Storti", "Poretti"),
  when = c("morning", "morning", "morning", "afternoon", "afternoon", "afternoon",
            "evening", "evening"),
  heartrate = c(67, 80, 64, 56, 90, 50, 60, 85)
)
```

It represents the heart rate measured on three patients in the morning, in the afternoon and in the evening. Make explicit any implicit missing value. How many missing values do you see?

```
heartrate %>% complete(surname, when, fill=list(name="Aldo"))
```

```
## # A tibble: 9 x 4
##   surname    when      name heartrate
##   <fctr>    <fctr>    <fctr>    <dbl>
## 1 Baglio afternoon  Aldo      56
## 2 Baglio evening    Aldo      NA
## 3 Baglio morning    Aldo      67
## 4 Poretti afternoon Giacomo    50
## 5 Poretti evening  Giacomo    85
## 6 Poretti morning  Giacomo    64
## 7 Storti afternoon Giovanni    90
## 8 Storti evening  Giovanni    60
## 9 Storti morning  Giovanni    80
```

```
# one missing value
```

```
# alternatively:
# heartrate %>%
# complete(surname, when) %>%
# fill(name)
```

7.1.2 Exercise 2

Import data in the file `marks.Rdata`. Missing values have been recorded as “.”. What’s the percentage of missing values in the data? Replace them with NA and drop them.

```
load("marks.Rdata")

marks_NA <- na_if(marks, ".")
marks_NA %>%
  filter(is.na(marks)) %>%
  summarise(n())/30

##           n()
## 1 0.06666667
```

```
# 6.7% of missing values
```

```
marks_NA %>% drop_na()
```

```
## # A tibble: 28 x 1
##   marks
##   <chr>
## 1    25
## 2    21
## 3    26
```

```
## 4    23
## 5    23
## 6    24
## 7    22
## 8    24
## 9    23
## 10   26
## # ... with 18 more rows
```

7.1.3 Exercise 3

Import the data `heartrate_NA.Rdta` . Consider all the missing values you find and replace them using the function `fill()` when possible.

```
load("heartrate_NA.Rdata")

heartrate_NA %>%
  na_if( "") %>%
  fill(name, surname)

## # A tibble: 9 x 4
##       name surname      when heartrate
##   <chr>   <chr>    <chr>      <dbl>
## 1    Aldo  Baglio  morning        67
## 2    Aldo  Baglio afternoon     56
## 3    Aldo  Baglio  evening        67
## 4 Giovanni Storti  morning        80
## 5 Giovanni Storti afternoon     90
## 6 Giovanni Storti  evening        60
## 7 Giacomo Poretti morning        64
## 8 Giacomo Poretti afternoon     50
## 9 Giacomo Poretti evening        85
```


Chapter 8

Dates with lubridate

8.1 lubridate

Note: all the exercises of this section are based on the `flights` dataset.

```
require(tidyverse)
require(lubridate)
require(nycflights13)
```

8.1.1 Exercise 1

Using the `flights` data, build the variable `dep_date` based on the variables `year`, `month` and `day`. First use the function `unite()` and then the parsing function `ymd()`. Select only the new variable and save the new data frame called `flights_date`.

```
flights_date <- flights %>%
  unite(date, year, month, day) %>%
  mutate(date = ymd(date)) %>%
  select(date)
```

8.1.2 Exercise 2

Using the dataset, shift all flights by two months. Save it in a separate data frame called `flights_date_2`.

```
flights_date_2 <- flights_date %>% mutate(date2 = date + months(2))
```

8.1.3 Exercise 3

Take the new date (2 months ahead) and subtract the original variable `date`. `flights_date_2`.

```
flights_date_2 %>% mutate(date2=date+2*month())

## # A tibble: 336,776 x 3
##       date       date2 `date2 - date`
##   <date>     <date>     <time>
## 1 2013-01-01 2013-03-01      59 days
## 2 2013-01-01 2013-03-01      59 days
## 3 2013-01-01 2013-03-01      59 days
## 4 2013-01-01 2013-03-01      59 days
## 5 2013-01-01 2013-03-01      59 days
## 6 2013-01-01 2013-03-01      59 days
## 7 2013-01-01 2013-03-01      59 days
## 8 2013-01-01 2013-03-01      59 days
## 9 2013-01-01 2013-03-01      59 days
## 10 2013-01-01 2013-03-01      59 days
## # ... with 336,766 more rows
```

Chapter 9

Manipulating strings with stringr

9.1 Data import: set working directory

In this section you will work with data are contained in the data folder. Hence you should set your working directory in the *data* folder, using `setwd()` function, like in this example:

```
setwd("~/Documents/datamanage/exercises/data)
```

You will work inside this folder.

```
library(tidyverse)
library(stringr)
```

9.1.1 Exercise 1

Import the data `aire_milano_strings.txt` which is a tab delimited file. Find how China has been codified (notice that the file is in Italian) and manipulate that string as you find more comfortable for you. Save the results in a new tibble.

```
## Parsed with column specification:
## cols(
##   Residenza = col_character(),
##   MotivoIscrizioneEsteri = col_character(),
##   Num = col_integer()
## )

aire %>% filter(str_detect(Residenza, c("Cina"))) # not recorded as Cina

## # A tibble: 0 x 3
## # ... with 3 variables: Residenza <chr>, MotivoIscrizioneEsteri <chr>,
## #   Num <int>
```

```

aire %>% filter(str_detect(Residenza, c("cina"))) # not recorded as cina

## # A tibble: 0 x 3
## # ... with 3 variables: Residenza <chr>, MotivoIscrizioneEsterio <chr>,
## #   Num <int>

aire %>% filter(str_detect(Residenza, c("Cin")))

## # A tibble: 5 x 3
##       Residenza      MotivoIscrizioneEsterio   Num
##       <chr>          <chr> <int>
## 1 Cinese, Rep. Popolare      all'emigrazione    535
## 2 Cinese, Rep. Popolare per acquisto cittadinanza    14
## 3 Cinese, Rep. Popolare      per nascita        119
## 4 Cinese, Rep. Popolare per residenza all'estero    26
## 5 Cinese, Rep. Popolare      trasferimento da AIRE    10

aire_clean <- aire %>% mutate(Residenza,
                             Residenza = str_replace(Residenza, c("Cinese, Rep. Popolare")
                                                         "Cina"))

```

9.1.2 Exercise 2

Using the data modified in exercise 1, find all the countries whose names contain non-alphanumeric characters. Identify what kind of characters they contain.

```

str_extract(aire_clean$Residenza, "[[:punct:]]")

```

9.1.3 Exercise 3

Consider now the column with information on the reason for migrating. Count how many different reasons there are and notice that citizenship was recorded in two slightly different ways: “acquisto cittadinanza” and “per acquisto cittadinanza”. Replace one of them so that they are the same.

```

# find all distinct reasons
aire_clean %>% distinct(MotivoIscrizioneEsterio)

## # A tibble: 8 x 1
##       MotivoIscrizioneEsterio
##       <chr>
## 1      all'emigrazione
## 2      per nascita
## 3 per residenza all'estero

```

```
## 4      acquisto cittadinanza
## 5          per matrimonio
## 6      trasferimento da AIRE
## 7 per acquisto cittadinanza
## 8          per sentenza
```

```
aire_clean <- aire_clean %>%
  mutate(MotivoIscrizioneEstero = str_replace(MotivoIscrizioneEstero,
                                              c("^acquisto cittadinanza"), "per acquisto cittadina
```

```
# now you only have 7 different levels
aire_clean %>% distinct(MotivoIscrizioneEstero)
```

```
## # A tibble: 7 x 1
##   MotivoIscrizioneEstero
##   <chr>
## 1 all'emigrazione
## 2 per nascita
## 3 per residenza all'estero
## 4 per acquisto cittadinanza
## 5 per matrimonio
## 6 trasferimento da AIRE
## 7 per sentenza
```


Chapter 10

Case study

10.1 Recap exercise

In this section you will use all the tools provided throughout the course. It is a simple example on real data that shows you the overall usage of the tools you were given during this course.

In the data folder you find the following three files:

1. `rating_final.csv`
2. `chefmezcuisine.csv`
3. `userprofile.csv`

These are the files you will work on in this section.

The files contain data on the rating of some restaurants in the US, their characteristics and the characteristics of the users rating them.

This exercise will guide you to first clean and combine data, and then answer questions on people's restaurant preferences.

Before starting the exercise, you should set your working directory in the *data* folder, using `setwd()` function, like in this example:

```
setwd("/data")
```

You will work in this folder.

10.2 Exercise

1. First of all you need to import the three files into R using the correct `readr` function. In order to find the correct function and to set the right options, you may explore the files by opening them in csv (check which is the separator, if there are column names, etc).

```

require(tidyverse)

# chefmozcuisine
chefmozcuisine <- read_delim("chefmozcuisine.csv", delim = ",", col_names = T)

## Parsed with column specification:
## cols(
##   placeID = col_integer(),
##   Rcuisine = col_character()
## )

# userprofile
userprofile <-
  read_delim("userprofile.csv", delim = ",", col_names = T)

## Parsed with column specification:
## cols(
##   userID = col_character(),
##   latitude = col_double(),
##   longitude = col_double(),
##   smoker = col_character(),
##   drink_level = col_character(),
##   dress_preference = col_character(),
##   ambience = col_character(),
##   transport = col_character(),
##   marital_status = col_character(),
##   hijos = col_character(),
##   birth_year = col_integer(),
##   interest = col_character(),
##   personality = col_character(),
##   religion = col_character(),
##   activity = col_character(),
##   color = col_character(),
##   weight = col_integer(),
##   budget = col_character(),
##   height = col_double()
## )

# rating_final
rating_final <-
  read_delim("rating_final.csv",
             delim = ",", col_names = T)

## Parsed with column specification:
## cols(
##   userID = col_character(),

```



```
## placeID = col_integer(),
## rating = col_integer(),
## food_rating = col_integer(),
## service_rating = col_integer()
## )
```

2. In order to understand what you are working on, check how many columns and rows each data frame is composed of, and check the type of the variables you are working with. If the variables type has not correctly been parsed, parse it manually. Further explore those variables that you think may be interesting for understanding different people's tastes, by for example checking how many levels they have, etc etc.

```
# chefmozcuisine
# how many rows? how many variables?
chefmozcuisine
```

```
## # A tibble: 916 x 2
##   placeID      Rcuisine
##   <int>      <chr>
## 1  135110    Spanish
## 2  135109    Italian
## 3  135107 Latin_American
## 4  135106    Mexican
## 5  135105    Fast_Food
## 6  135104    Mexican
## 7  135103    Burgers
## 8  135103 Dessert-Ice_Cream
## 9  135103    Fast_Food
## 10 135103    Hot_Dogs
## # ... with 906 more rows
```

```
# how many different types of cuisine?
chefmozcuisine %>% distinct(Rcuisine)
```

```
## # A tibble: 59 x 1
##   Rcuisine
##   <chr>
## 1    Spanish
## 2    Italian
## 3 Latin_American
## 4    Mexican
## 5    Fast_Food
## 6    Burgers
## 7 Dessert-Ice_Cream
## 8    Hot_Dogs
## 9    Steaks
```

```
## 10           Asian
## # ... with 49 more rows

# userprofile
# how many rows? how many variables?
userprofile

## # A tibble: 138 x 19
##   userID latitude longitude smoker   drink_level dress_preference ambience
##   <chr>    <dbl>    <dbl> <chr>      <chr>          <chr>    <chr>
## 1 U1001 22.14000 -100.9788 false    abstemious    informal  family
## 2 U1002 22.15009 -100.9833 false    abstemious    informal  family
## 3 U1003 22.11985 -100.9465 false    social drinker    formal    family
## 4 U1004 18.86700 -99.1830 false    abstemious    informal  family
## 5 U1005 22.18348 -100.9599 false    abstemious    no preference family
## 6 U1006 22.15000 -100.9830 true     social drinker    no preference friends
## 7 U1007 22.11846 -100.9383 false    casual drinker    informal  solitary
## 8 U1008 22.12299 -100.9238 false    social drinker    formal    solitary
## 9 U1009 22.15943 -100.9904 false    abstemious    formal    family
## 10 U1010 22.19089 -100.9987 false    social drinker    no preference friends
## # ... with 128 more rows, and 12 more variables: transport <chr>,
## #   marital_status <chr>, hijos <chr>, birth_year <int>, interest <chr>,
## #   personality <chr>, religion <chr>, activity <chr>, color <chr>,
## #   weight <int>, budget <chr>, height <dbl>

# let us explore some interesting variables that may be interesting in future analysis
userprofile %>% distinct(activity)

## # A tibble: 5 x 1
##   activity
##   <chr>
## 1 student
## 2 professional
## 3 ?
## 4 unemployed
## 5 working-class

userprofile %>% distinct(ambience)

## # A tibble: 4 x 1
##   ambience
##   <chr>
## 1 family
## 2 friends
## 3 solitary
```

```
## 4      ?
```

```
userprofile %>% distinct(smoker)
```

```
## # A tibble: 3 x 1
##   smoker
##   <chr>
## 1 false
## 2  true
## 3    ?
```

```
# rating_final
# how many rows?
rating_final
```

```
## # A tibble: 1,161 x 5
##   userID placeID rating food_rating service_rating
##   <chr>   <int> <int>      <int>      <int>
## 1 U1077  135085     2         2         2
## 2 U1077  135038     2         2         1
## 3 U1077  132825     2         2         2
## 4 U1077  135060     1         2         2
## 5 U1068  135104     1         1         2
## 6 U1068  132740     0         0         0
## 7 U1068  132663     1         1         1
## 8 U1068  132732     0         0         0
## 9 U1068  132630     1         1         1
## 10 U1067  132584     2         2         2
## # ... with 1,151 more rows
```

```
rating_final %>% distinct(rating)
```

```
## # A tibble: 3 x 1
##   rating
##   <int>
## 1     2
## 2     1
## 3     0
```

```
rating_final %>% distinct(food_rating)
```

```
## # A tibble: 3 x 1
##   food_rating
```

```
##           <int>
## 1           2
## 2           1
## 3           0
```

```
rating_final %>% distinct(service_rating)
```

```
## # A tibble: 3 x 1
##   service_rating
##           <int>
## 1             2
## 2             1
## 3             0
```

```
# ratings are either 0, 1 or 2
```

```
# notice that placeID is an integer value. However, it is the ID hence you will
# calculate no statistics on it. You may as well force it to a character vector.
```

```
rating_final <- rating_final %>%
  mutate(placeID = as.character(placeID))
```

```
chefmozcuisine <- chefmozcuisine %>%
  mutate(placeID = as.character(placeID))
```

3. Consider `userprofile` data frame. Create a new data frame containing only relevant variables for users profile. Among these, keep the variables: `userID`, `birth_year`, `budget`, `marital_status`, `personality`, `smoker` and `activity`. If you think there are other relevant variables, you may include them in the new data frame as well. Call the new data frame `userprofile_reduced`.

```
userprofile_reduced <- userprofile %>%
  select(userID, birth_year, budget, marital_status, personality, smoker, activity)
```

```
userprofile_reduced
```

```
## # A tibble: 138 x 7
##   userID birth_year budget marital_status personality smoker
##   <chr>    <int> <chr>         <chr>         <chr> <chr>
## 1 U1001    1989 medium      single  thrifty-protector false
## 2 U1002    1990 low        single  hunter-ostentatious false
## 3 U1003    1989 low        single  hard-worker false
## 4 U1004    1940 medium     single  hard-worker false
## 5 U1005    1992 medium     single  thrifty-protector false
## 6 U1006    1989 medium     single  hard-worker true
## 7 U1007    1989 low        single  thrifty-protector false
```

```
## 8 U1008      1989    low      single      hard-worker false
## 9 U1009      1991 medium    single    thrifty-protector false
## 10 U1010     1987 medium    married      hard-worker false
## # ... with 128 more rows, and 1 more variables: activity <chr>
```

4. Focus on the data frame `userprofile_reduced`. By exploring the different values recorded for `budget`, you may notice there are missing values. What are they recorded by? Replace all missing values with `NA`. Do the same for all the variables in `userprofile_reduced`.

```
userprofile_reduced %>%
  distinct(budget)
```

```
## # A tibble: 4 x 1
##   budget
##   <chr>
## 1 medium
## 2    low
## 3     ?
## 4   high
```

"?" is for missing values

```
userprofile_reduced <- userprofile_reduced %>% na_if("?")
```

5. Note that for all users we have the year of birth but we do not have the age. As it may be easier to deal with their age, build a variable called `age` and drop `birth_year`.

```
require(lubridate)
```

```
## Loading required package: lubridate
```

```
## Loading required package: methods
```

```
##
## Attaching package: 'lubridate'
```

```
## The following object is masked from 'package:base':
##
##   date
```

```

userprofile_reduced <- userprofile_reduced %>%
  mutate(age = year(today())-birth_year) %>%
  select(-birth_year)

```

6. All three data frames are now ready to use. Merge the three data frames (`rating_final`, `userprofile_reduced` and `chefmoicuisine`) so that you keep all rows and columns of rating and you add all the variables of `chefmez cuisine` and `userprofile_reduced`. Call the new data frame `rating_all`.

```

rating_all <- rating_final %>%
  left_join(userprofile_reduced) %>%
  left_join(chefmoz cuisine)

```

```
## Joining, by = "userID"
```

```
## Joining, by = "placeID"
```

7. Group data by `placeID` and find the average rating (`rating`, `food_rating` and `service_rating`). Sort the results so that places with the highest average rating are at the top. Show id of such places and type of cuisine. What is the cuisine type of users' favourite restaurants?

```

rating_all %>%
  group_by(placeID) %>%
  summarise_at(vars(rating_mean = rating, food_rating_mean = food_rating,
                    service_rating_mean = service_rating),
               funs(mean)) %>%
  arrange(desc(rating_mean), desc(food_rating_mean), desc(service_rating_mean)) %>%
  left_join(chefmoz cuisine)

```

```
## Joining, by = "placeID"
```

```
## # A tibble: 147 x 5
##   placeID rating_mean food_rating_mean service_rating_mean Rcuisine
##   <chr>      <dbl>          <dbl>          <dbl>          <chr>
## 1 134986    2.000000          2.00          2.000000 International
## 2 135034    2.000000          2.00          1.600000 Japanese
## 3 132955    2.000000          1.80          1.800000 Bar_Pub_Brewery
## 4 132922    1.833333          1.50          1.833333 Cafeteria
## 5 132755    1.800000          2.00          1.600000 Mexican
## 6 135013    1.750000          2.00          1.750000 <NA>
## 7 135074    1.750000          1.75          1.750000 Contemporary
## 8 134976    1.750000          1.75          1.000000 Mexican
## 9 134976    1.750000          1.75          1.000000 Mediterranean
## 10 134976    1.750000          1.75          1.000000 Burgers
## # ... with 137 more rows
```

8. Find mean and standard deviation of all rating variables. Do you notice differences with regards to ratings of students as compared to people that are employed? Do you find differences in smokers and non smokers? Think of other groups of users to compare.

```
rating_all %>%
  group_by(activity) %>%
  summarise_at(vars(rating, food_rating, service_rating),
    funs(n(), mean, sd))

## # A tibble: 5 x 10
##   activity rating_n food_rating_n service_rating_n rating_mean
##   <chr>      <int>      <int>          <int>      <dbl>
## 1 professional    135        135            135    1.385185
## 2 student      1114       1114           1114    1.188510
## 3 unemployed     15         15             15    0.000000
## 4 working-class   4          4              4    1.500000
## 5 <NA>          63         63             63    1.365079
## # ... with 5 more variables: food_rating_mean <dbl>, service_rating_mean <dbl>,
## #   rating_sd <dbl>, food_rating_sd <dbl>, service_rating_sd <dbl>

rating_all %>%
  group_by(smoker) %>%
  summarise_at(vars(rating, food_rating, service_rating),
    funs(n(), mean, sd))

## # A tibble: 3 x 10
##   smoker rating_n food_rating_n service_rating_n rating_mean food_rating_mean
##   <chr>      <int>      <int>          <int>      <dbl>      <dbl>
## 1 false    1034        1034            1034    1.210832    1.223404
## 2 true      259         259             259    1.150579    1.108108
## 3 <NA>       38          38              38    1.394737    1.315789
## # ... with 4 more variables: service_rating_mean <dbl>, rating_sd <dbl>,
## #   food_rating_sd <dbl>, service_rating_sd <dbl>

rating_all %>%
  group_by(budget) %>%
  summarise_at(vars(rating, food_rating, service_rating),
    funs(n(), mean, sd))

## # A tibble: 4 x 10
##   budget rating_n food_rating_n service_rating_n rating_mean food_rating_mean
##   <chr>      <int>      <int>          <int>      <dbl>      <dbl>
## 1 high      49         49             49    1.469388    1.612245
## 2 low       362        362            362    1.127072    1.096685
## 3 medium    838        838            838    1.219570    1.226730
## 4 <NA>      82         82             82    1.231707    1.195122
```

```
## # ... with 4 more variables: service_rating_mean <dbl>, rating_sd <dbl>,
## #   food_rating_sd <dbl>, service_rating_sd <dbl>
```

9. What are the best restaurant for smokers? What are the best restaurant for those that are on a budget?

```
rating_all %>%
  filter(smoker == "true") %>%
  group_by(placeID) %>%
  summarise_at(vars(rating_mean = rating, food_rating_mean = food_rating,
                    service_rating_mean = service_rating),
               funs(mean, n())) %>%
  arrange(desc(rating_mean), desc(food_rating_mean), desc(service_rating_mean)) %>%
  left_join(chefmozcuisine)
```

```
## Joining, by = "placeID"
```

```
## # A tibble: 120 x 8
##   placeID rating_mean food_rating_mean service_rating_mean rating_n
##   <chr>      <dbl>          <dbl>          <dbl>      <int>
## 1 132660      2            2            2          1
## 2 132755      2            2            2          1
## 3 132767      2            2            2          1
## 4 132830      2            2            2          1
## 5 132884      2            2            2          1
## 6 132958      2            2            2          1
## 7 134986      2            2            2          2
## 8 134999      2            2            2          1
## 9 135000      2            2            2          1
## 10 135019      2            2            2          1
## # ... with 110 more rows, and 3 more variables: food_rating_n <int>,
## #   service_rating_n <int>, Rcuisine <chr>
```

```
rating_all %>%
  filter(budget == "low") %>%
  group_by(placeID) %>%
  summarise_at(vars(rating_mean = rating, food_rating_mean = food_rating,
                    service_rating_mean = service_rating),
               funs(mean, n())) %>%
  arrange(desc(rating_mean), desc(food_rating_mean), desc(service_rating_mean)) %>%
  left_join(chefmozcuisine)
```

```
## Joining, by = "placeID"
```



```
## # A tibble: 120 x 8
##   placeID rating_mean food_rating_mean service_rating_mean rating_n
##   <chr>      <dbl>          <dbl>          <dbl>          <int>
## 1 132884      2            2.00            2.0            2
## 2 132958      2            2.00            2.0            1
## 3 134975      2            2.00            2.0            1
## 4 134996      2            2.00            2.0            1
## 5 135049      2            2.00            2.0            1
## 6 135055      2            2.00            2.0            2
## 7 132608      2            2.00            1.5            2
## 8 132723      2            2.00            1.5            2
## 9 132755      2            2.00            1.5            2
## 10 135075      2            1.75            1.5            4
## # ... with 110 more rows, and 3 more variables: food_rating_n <int>,
## #   service_rating_n <int>, Rcuisine <chr>

rating_all %>%
  filter(budget == "high") %>%
  group_by(placeID) %>%
  summarise_at(vars(rating_mean = rating, food_rating_mean = food_rating,
                    service_rating_mean = service_rating),
               funs(mean, n())) %>%
  arrange(desc(rating_mean), desc(food_rating_mean), desc(service_rating_mean)) %>%
  left_join(chefmozcuisine)

## Joining, by = "placeID"

## # A tibble: 37 x 8
##   placeID rating_mean food_rating_mean service_rating_mean rating_n
##   <chr>      <dbl>          <dbl>          <dbl>          <int>
## 1 132584      2            2            2            1
## 2 132733      2            2            2            1
## 3 132862      2            2            2            1
## 4 134986      2            2            2            1
## 5 134996      2            2            2            1
## 6 135025      2            2            2            1
## 7 135028      2            2            2            1
## 8 135030      2            2            2            2
## 9 135030      2            2            2            2
## 10 135039      2            2            2            1
## # ... with 27 more rows, and 3 more variables: food_rating_n <int>,
## #   service_rating_n <int>, Rcuisine <chr>
```

notice however that each restaurant has only few ratings (often just 1!)

10. Think of other analysis that you may perform with these data and use the tidyverse toolbox to answer your questions!